

Docker-Kubernetes-Foundations

Linux Namespaces

Exercise-1

PID Namespaces

NOTE: Perform the below steps on the **dnodel** node

1. PID of current bash process

- Login to a new terminal in **dnodel** node and verify the PID of current bash process in global NS

```
ps $$
```

```
root@dnodel:~# ps $$
  PID TTY          STAT       TIME COMMAND
 21441 pts/0        Ss           0:00 -bash
root@dnodel:~#
```

- Read the PID namespace of the bash process

```
readlink /proc/$$/ns/pid
```

```
root@dnodel:~# readlink /proc/$$/ns/pid
pid:[4026531836]
root@dnodel:~#
```

2. New PID namespace for a bash process

- Enter the superuser mode

```
sudo su
```

- Create a new PID namespace for a bash process

```
unshare -pf --mount-proc /bin/bash
```

- Read the PID namespace of the bash process

```
readlink /proc/$$/ns/pid
```

```
root@dnode1:~# sudo su
root@dnode1:/home/ubuntu# unshare -pf --mount-proc /bin/bash
root@dnode1:/home/ubuntu# readlink /proc/$$/ns/pid
pid:[4026532197]
root@dnode1:/home/ubuntu#
```

3. Running processes

- Verify the current running processes in new PID NS

```
ps -ef
```

```
root@dnode1:/home/ubuntu# ps -ef
UID          PID    PPID  C STIME TTY          TIME CMD
root           1         0  0  05:56 pts/0        00:00:00 /bin/bash
root          10         1  0  05:57 pts/0        00:00:00 ps -ef
root@dnode1:/home/ubuntu#
```

- Create new process in this new bash

```
sleep 10000 &
```

```
root@dnode1:/home/ubuntu# sleep 10000 &
[1] 11
root@dnode1:/home/ubuntu#
```

- Verify the current running processes in new PID NS

```
ps -ef
```

```
root@dnode1:/home/ubuntu# ps -ef
UID          PID    PPID  C STIME TTY          TIME CMD
root           1         0  0  05:56 pts/0        00:00:00 /bin/bash
root          11         1  0  05:58 pts/0        00:00:00 sleep 10000
root          12         1  0  05:58 pts/0        00:00:00 ps -ef
root@dnode1:/home/ubuntu#
```

Observe the Parent and Child Process ID in new PID NS

- Open another SSH session for **dnode1**, verify the current running processes in global PID NS

```
ps -ef
```

Criterion Networks

```
daemon 1079 1 0 Jul21 ? 00:00:00 /usr/sbin/atd -f
root 1084 1 0 Jul21 ? 00:00:00 /usr/bin/lxcfs /var/lib/lxcfs/
root 1091 1 0 Jul21 ? 00:00:04 /lib/systemd/systemd-logind
root 1137 1 0 Jul21 ? 00:00:00 /usr/sbin/cron -f
syslog 1177 1 0 Jul21 ? 00:00:01 /usr/sbin/rsyslogd -n
message+ 1178 1 0 Jul21 ? 00:00:07 /usr/bin/dbus-daemon --system --address=systemd: --nofork --nospidfile --systemd-activation
root 1216 1 0 Jul21 ? 00:00:00 /usr/sbin/acpid
root 1218 1 0 Jul21 ? 00:00:03 /usr/lib/accounts-service/accounts-daemon
root 1221 1 0 Jul21 ? 00:00:06 /usr/lib/snapd/snapd
root 1261 1 0 Jul21 ? 00:00:00 /sbin/mdadm --monitor --pid-file /run/mdadm/monitor.pid --daemonise --scan --syslog
root 1287 1 0 Jul21 ? 00:00:01 /usr/lib/policykit-1/polkitd --no-debug
root 1312 1 0 Jul21 tty1 00:00:00 /sbin/agetty --noclear tty1 linux
root 1354 1 0 Jul21 ? 00:00:02 /usr/sbin/irqbalance --pid=/var/run/irqbalance.pid
root 12168 1 0 01:12 tty80 00:00:00 /sbin/agetty --keep-baud 115200 38400 9600 tty80 vt220
root 17095 2 0 04:35 ? 00:00:00 [kworker/u16:2]
root 17455 2 0 04:43 ? 00:00:00 [kworker/u16:0]
root 18234 1 0 04:56 ? 00:00:00 /usr/sbin/sshd -D
root 19452 2 0 05:07 ? 00:00:00 [kworker/1:4]
root 19756 2 0 05:12 ? 00:00:00 [kworker/0:1]
root 20847 2 0 05:22 ? 00:00:00 [kworker/u16:3]
root 21206 2 0 05:51 ? 00:00:00 [kworker/1:0]
root 21294 2 0 05:51 ? 00:00:00 [kworker/0:2]
root 21359 18234 0 05:54 ? 00:00:00 sshd: ubuntu@pts/0
root 21361 1 0 05:54 ? 00:00:00 /lib/systemd/systemd --user
root 21362 21361 0 05:54 ? 00:00:00 (sd-pam)
root 21441 21359 0 05:54 pts/0 00:00:00 -bash
root 21466 21441 0 05:56 pts/0 00:00:00 sudo su
root 21468 21466 0 05:56 pts/0 00:00:00 su
root 21471 21468 0 05:56 pts/0 00:00:00 bash
root 21479 21471 0 05:56 pts/0 00:00:00 unshare -pf --mount-proc /bin/bash
root 21480 21479 0 05:56 pts/0 00:00:00 /bin/bash
root 21490 21480 0 05:58 pts/0 00:00:00 sleep 10000
root 21492 18234 0 05:59 ? 00:00:00 sshd: ubuntu@pts/1
root 21535 21492 0 05:59 pts/1 00:00:00 -bash
root 21552 21535 0 05:59 pts/1 00:00:00 ps -ef
root@dnodel1:~#
```

Observe the Process ID for bash and sleep process in global and new PID NS

Exercise-2

UTS namespace

NOTE: Perform the below steps on the *dnodel* node

1. PID of current bash process

- Login to a new terminal in **dnodel** node and verify the PID of current bash process in global NS

```
ps $$
```

```
Ubuntu 16.04.4 LTS
Welcome to Ubuntu 16.04.4 LTS (GNU/Linux 4.14.0-041400-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

Get cloud support with Ubuntu Advantage Cloud Guest:
http://www.ubuntu.com/business/services/cloud

54 packages can be updated.
1 update is a security update.

Last login: Sun Jul 22 05:59:13 2018 from 192.168.122.1
root@dnodel:~# ps $$
  PID TTY          STAT       TIME COMMAND
 21656 pts/0        Ss           0:00 -bash
root@dnodel:~#
```

- Read the PID namespace of the bash process

```
readlink /proc/$$/ns/pid
```

```
root@dnodel:~# readlink /proc/$$/ns/pid
pid:[4026531836]
root@dnodel:~#
```

- Read the UTS namespace of the bash process

```
readlink /proc/$$/ns/uts
```

```
root@dnode1:~# readlink /proc/$$/ns/uts
uts:[4026531838]
root@dnode1:~#
```

- Verify the current hostname of the device

```
hostname
```

2. New UTS namespace for a bash process

- Create a new UTS namespace for a bash process

```
unshare -u /bin/bash
```

- Read the PID namespace of the bash process

```
readlink /proc/$$/ns/pid
```

```
root@dnode1:~# unshare -u /bin/bash
root@dnode1:~# readlink /proc/$$/ns/pid
pid:[4026531836]
root@dnode1:~#
```

- Verify the new bash process PID

```
ps $$
```

- Read the UTS namespace of the bash process

```
readlink /proc/$$/ns/uts
```

```
root@dnode1:~# readlink /proc/$$/ns/uts
uts:[4026532196]
root@dnode1:~#
```

- Verify that UTS inode is different but PID inode is same when compared to global namespace
- Change the hostname in this new UTS shell

```
hostname helloworld
```

```
root@dnode1:~# hostname helloworld
root@dnode1:~# hostname
helloworld
root@dnode1:~#
```

- Display the changed hostname

```
hostname
```

- Exit the UTS namespace

```
exit
```

- Verify the hostname of the device is unchanged

```
hostname
```

```
root@dnode1:~# exit
exit
root@dnode1:~# hostname
dnode1
root@dnode1:~#
```

Exercise-3

User Namespace

NOTE: Perform the below steps on the *dnode1* node

Create a new user

- Login to a new terminal in *dnode1* node and create a new user named *john*

```
useradd john
```

```
Ubuntu 16.04.4 LTS
Welcome to Ubuntu 16.04.4 LTS (GNU/Linux 4.14.0-041400-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage

Get cloud support with Ubuntu Advantage Cloud Guest:
http://www.ubuntu.com/business/services/cloud

54 packages can be updated.
1 update is a security update.

Last login: Sun Jul 22 06:00:55 2018 from 192.168.122.1
root@dnode1:~# useradd john
root@dnode1:~#
```

- Assign Password to *john*

```
passwd john
```

```
root@dnode1:~# passwd john
Enter new UNIX password:
Retype new UNIX password:
passwd: password updated successfully
root@dnode1:~#
```

- Switch user to *john*

```
su john
```

```
root@dnodel:~# su john  
john@dnodel:/home/ubuntu$
```

- Verify the user for current process

```
whoami
```

```
john@dnodel:/home/ubuntu$ whoami  
john  
john@dnodel:/home/ubuntu$
```

Create new User namespace

- Create new User namespace

```
unshare --map-root-user -U /bin/bash
```

```
john@dnodel:/home/ubuntu$ unshare --map-root-user -U /bin/bash  
root@dnodel:/home/ubuntu# whoami  
root  
root@dnodel:/home/ubuntu#
```

- Verify the user for current process

```
whoami
```

Verify the output is root

Exercise-4

Mount Namespaces

NOTE: Perform the below steps on the *dnode1* node

Create a new mount space, mount point and verify them

- Enter the superuser mode

```
sudo su
```

```
Ubuntu 16.04.4 LTS
Welcome to Ubuntu 16.04.4 LTS (GNU/Linux 4.14.0-041400-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage

Get cloud support with Ubuntu Advantage Cloud Guest:
http://www.ubuntu.com/business/services/cloud

54 packages can be updated.
1 update is a security update.

Last login: Sun Jul 22 06:10:39 2018 from 192.168.122.1
root@dnode1:~# sudo su
root@dnode1:/home/ubuntu#
```

- Login to a new terminal in *dnode1* node and create new mount namespace

```
unshare -m /bin/bash
```

- Create a new directory. Please observe that a new directory is created under */tmp* with random ID.

```
secret_dir=`mktemp -d --tmpdir=/tmp`
```

- Creating a new mount point for above created directory

```
mount -n -o size=1m -t tmpfs tmpfs $secret_dir
```

```
root@dnode1:/home/ubuntu# unshare -m /bin/bash
root@dnode1:/home/ubuntu# secret_dir=`mktemp -d --tmpdir=/tmp`
root@dnode1:/home/ubuntu# mount -n -o size=1m -t tmpfs tmpfs $secret_dir
```

Criterion Networks

- Check the available mount points

```
grep /tmp /proc/mounts
```

```
root@dnode1:/home/ubuntu# grep /tmp /proc/mounts
tmpfs /tmp/tmp.JmpVLwssUm tmpfs rw,relatime,size=1024k 0 0
root@dnode1:/home/ubuntu# cd /tmp/tmp.JmpVLwssUm
root@dnode1:/tmp/tmp.JmpVLwssUm#
```

- Change to the new directory (replace tmpID with actual directory name created above) and create 2 files in new directory

```
cd /tmp/tmpID
```

```
touch hello
```

```
touch bye
```

```
root@dnode1:/tmp/tmp.JmpVLwssUm# touch hello
root@dnode1:/tmp/tmp.JmpVLwssUm# touch bye
root@dnode1:/tmp/tmp.JmpVLwssUm#
```

- List the files created

```
ls -al
```

```
root@dnode1:/tmp/tmp.JmpVLwssUm# ls -al
total 4
drwxrwxrwt 2 root root 80 Jul 22 06:17 .
drwxrwxrwt 9 root root 4096 Jul 22 06:17 ..
-rw-r--r-- 1 root root 0 Jul 22 06:17 bye
-rw-r--r-- 1 root root 0 Jul 22 06:17 hello
root@dnode1:/tmp/tmp.JmpVLwssUm#
```

- Open another session to dnode1 and change directory within /tmp to newly created directory
- List the files

```
ls -al
```

Verify we cannot see any files that were created earlier

Exercise-5

Net Namespaces

NOTE: Perform the below steps on the *dnode1* node

- Login to a new terminal in *dnode1* node and verify the current network interfaces

```
ip link list
```

```
Ubuntu 16.04.4 LTS
Welcome to Ubuntu 16.04.4 LTS (GNU/Linux 4.14.0-041400-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

Get cloud support with Ubuntu Advantage Cloud Guest:
http://www.ubuntu.com/business/services/cloud

54 packages can be updated.
1 update is a security update.

Last login: Sun Jul 22 06:13:41 2018 from 192.168.122.1
root@dnode1:~# ip link list
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: ens3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT group default qlen 1000
    link/ether 52:54:00:db:01:78 brd ff:ff:ff:ff:ff:ff
3: ens4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT group default qlen 1000
    link/ether 52:54:00:04:8b:94 brd ff:ff:ff:ff:ff:ff
root@dnode1:~#
```

- Create a new net namespace

```
ip netns add tom
```

- Verify the namespace created

```
ip netns list
```

```
root@dnode1:~# ip netns add tom
root@dnode1:~# ip netns list
tom
root@dnode1:~#
```

- Create a virtual Ethernet pair

```
ip link add veth100 type veth peer name veth101
```

```
ip link list
```

```
ifconfig veth100 up
```

```
ifconfig veth101 up
```

```
root@dnode1:~# ip link add veth100 type veth peer name veth101
root@dnode1:~# ip link list
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: ens3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT group default qlen 1000
   link/ether 52:54:00:db:01:78 brd ff:ff:ff:ff:ff:ff
3: ens4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP mode DEFAULT group default qlen 1000
   link/ether 52:54:00:04:8b:94 brd ff:ff:ff:ff:ff:ff
4: veth101@veth100: <BROADCAST,MULTICAST,M-DOWN> mtu 1500 qdisc noop state DOWN mode DEFAULT group default qlen 1000
   link/ether 8e:ef:ae:ca:a2:59 brd ff:ff:ff:ff:ff:ff
5: veth100@veth101: <BROADCAST,MULTICAST,M-DOWN> mtu 1500 qdisc noop state DOWN mode DEFAULT group default qlen 1000
   link/ether ee:5e:eb:bb:2b:7f brd ff:ff:ff:ff:ff:ff
root@dnode1:~# ifconfig veth100 up
root@dnode1:~# ifconfig veth101 up
root@dnode1:~#
```

- Move the veth101 interface into tom namespace

```
ip link set veth101 netns tom
```

- Verify the interfaces in tom namespace

```
ip netns exec tom ip link list
```

- Assign IP address to veth101 and make it up

```
ip netns exec tom ifconfig veth101 3.3.3.2/24 up
```

```
root@dnode1:~# ip netns exec tom ip link list
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN mode DEFAULT group default qlen 1000
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
4: veth101@if5: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN mode DEFAULT group default qlen 1000
   link/ether 8e:ef:ae:ca:a2:59 brd ff:ff:ff:ff:ff:ff link-netnsid 0
root@dnode1:~# ip netns exec tom ifconfig veth101 3.3.3.2/24 up
root@dnode1:~#
```

- Install bridge-utils package

```
apt-get install bridge-utils
```

- Create a new linux bridge br100 and make it up

```
brctl addbr br100
```

```
ifconfig br100 up
```

```

root@dnode1:~# apt-get install bridge-utils
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  bridge-utils
0 upgraded, 1 newly installed, 0 to remove and 53 not upgraded.
Need to get 28.6 kB of archives.
After this operation, 102 kB of additional disk space will be used.
0% [Working]
Get:1 http://archive.ubuntu.com/ubuntu xenial/main amd64 bridge-utils amd64
Fetched 28.6 kB in 0s (86.4 kB/s)
Selecting previously unselected package bridge-utils.
(Reading database ... 117008 files and directories currently installed.)
Preparing to unpack .../bridge-utils_1.5-9ubuntu1_amd64.deb ...
Unpacking bridge-utils (1.5-9ubuntu1) ...
Processing triggers for man-db (2.7.5-1) ...
Setting up bridge-utils (1.5-9ubuntu1) ...
root@dnode1:~# brctl addbr br100
root@dnode1:~# ifconfig br100 up
root@dnode1:~#

```

- Attach veth100 interface to bridge br100

```
brctl addif br100 veth100
```

- Assign IP address to interface br100 in global namespace

```
ifconfig br100 3.3.3.1/24 up
```

- Verify the network connectivity from tom and global namespace

```
ip netns exec tom ping 3.3.3.1
```

```

root@dnode1:~# brctl addif br100 veth100
root@dnode1:~# ifconfig br100 3.3.3.1/24 up
root@dnode1:~# ip netns exec tom ping 3.3.3.1
PING 3.3.3.1 (3.3.3.1) 56(84) bytes of data.
64 bytes from 3.3.3.1: icmp_seq=1 ttl=64 time=0.050 ms
64 bytes from 3.3.3.1: icmp_seq=2 ttl=64 time=0.038 ms
64 bytes from 3.3.3.1: icmp_seq=3 ttl=64 time=0.046 ms
64 bytes from 3.3.3.1: icmp_seq=4 ttl=64 time=0.042 ms

--- 3.3.3.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3061ms
rtt min/avg/max/mdev = 0.038/0.044/0.050/0.004 ms
^Croot@dnode1:~#

```

Exercise-6

Multiple Namespaces (PID and UTS)

NOTE: *Perform the below steps on dnode1*

- Enter the superuser mode

```
sudo su
```

- Create bash process with multiple namespaces

```
unshare -fp --mount-proc --pid --uts /bin/bash
```

- Verify the PID and UTS namespaces associated with this process ((Refer to previous exercises for commands))
- Exit from this namespace
- Let us try namespace within namespace for PID

```
unshare -fp --mount-proc /bin/bash
```

```
unshare -fp --mount-proc /bin/bash
```

```
unshare -fp --mount-proc /bin/bash
```

```
^Croot@dnode1:~# sudo su
root@dnode1:/home/ubuntu# unshare -fp --mount-proc --pid --uts /bin/bash
root@dnode1:/home/ubuntu# unshare -fp --mount-proc /bin/bash
root@dnode1:/home/ubuntu# unshare -fp --mount-proc /bin/bash
root@dnode1:/home/ubuntu# unshare -fp --mount-proc /bin/bash
root@dnode1:/home/ubuntu#
```

- Verify the total processes created in global namespace

```
ps -ef | grep bash
```

- Observe child parent-child relationship among 3 levels of namespaces within them

```
pstree
```

```
root@dnode1:/home/ubuntu# ps -ef | grep bash
root      1      0  0 06:36 pts/1    00:00:00 /bin/bash
root     10      1  0 06:37 pts/1    00:00:00 grep --color=auto bash
root@dnode1:/home/ubuntu# pstree
bash--pstree
root@dnode1:/home/ubuntu#
```

Linux cgroups

Exercise-7

Explore command to limit cpu usage of process

NOTE: *Perform the below exercises on "dnode1" through SSH and also on "desktop1" through VNC*

- Install prime number generator for generating huge cpu usage

```
apt-get update  
apt install libmath-prime-util-gmp-perl
```

- Verify the primes command between 1 100

```
primes 1 100
```

- Generate prime number for large set and redirect to null

```
primes 0 999999999 > /dev/null &
```

- Open VNC session to the "desktop1" machine. Please check Cluster Access -> VNC Access to get more details.
- Open a terminal within VNC desktop1 machine. It will be located at the top of the VNC session.
- Get the ens3 interface IP of the **dnode1**

```
ifconfig ens3
```

- SSH from the terminal opened in VNC session.

```
ssh ubuntu@192.168.122.205
```

Password: ubuntu

- On "dnode1" terminal (opened in VNC terminal), and switch to root user:

```
sudo su
```

- Issue top command to see the CPU usage of processes

```
top
```


Criterion Networks

- Verify that perl/primes command is consuming ~100% CPU
- Keep the VNC session to the "desktop1" machine running. We will use it for later exercises.
- On dnode1, bring the process into foreground

```
fg
```

(and then press ctrl +c)

CONFIDENTIAL

Exercise-8

Use `cpulimit` command to restrict the CPU usage of a process

NOTE: *Perform below exercise on dnode1*

- Install `cpulimit` command

```
apt-get install cpulimit
```

```
root@dnode1:~# apt-get install cpulimit
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  cpulimit
0 upgraded, 1 newly installed, 0 to remove and 133 not upgraded.
Need to get 15.1 kB of archives.
After this operation, 64.5 kB of additional disk space will be used.
Get:1 http://archive.ubuntu.com/ubuntu xenial/universe amd64 cpulimit amd64 2.2-1 [15.1 kB]
Fetched 15.1 kB in 1s (11.0 kB/s)
Selecting previously unselected package cpulimit.
(Reading database ... 117203 files and directories currently installed.)
Preparing to unpack .../cpulimit_2.2-1_amd64.deb ...
Unpacking cpulimit (2.2-1) ...
Processing triggers for man-db (2.7.5-1) ...
Setting up cpulimit (2.2-1) ...
root@dnode1:~#
```

- Execute `primes` command with `cpulimit` at the beginning

```
cpulimit -l 50 primes 0 9999999999 > /dev/null &
```

```
root@dnode1:~# cpulimit -l 50 primes 0 9999999999 > /dev/null &
[1] 22761
root@dnode1:~#
```

- From VNC, observe from `top` command that process is limited to 50% of CPU

desktop-kube-test (root's X desktop (ubuntu:1)) - VNC Viewer

Applications Home root@dnodel: /home/ubuntu 06:01 Ubuntu

root@dnodel: /home/ubuntu

File Edit View Search Terminal Help

```
top - 11:23:26 up 1 day, 59 min, 2 users, load average: 0.25, 0.15, 0.09
Tasks: 96 total, 1 running, 54 sleeping, 1 stopped, 0 zombie
%Cpu(s): 23.7 us, 0.2 sy, 0.0 ni, 73.1 id, 0.0 wa, 0.0 hi, 0.0 st, 3.0 st
KiB Mem : 2043264 total, 1374056 free, 93168 used, 576040 buff/cache
KiB Swap: 0 total, 0 free, 0 used. 1719192 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
22762	root	20	0	61560	38864	5000	T	51.5	1.9	0:11.16	perl
8	root	20	0	0	0	0	I	2.3	0.0	0:04.66	rcu_sched
22380	root	20	0	92788	6540	5616	S	0.3	0.3	0:00.13	sshd
1	root	20	0	37876	5828	3916	S	0.0	0.3	0:18.03	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.01	kthreadd
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/0:+
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	mm_percpu_+
7	root	20	0	0	0	0	S	0.0	0.0	0:00.48	ksoftirqd/0
9	root	20	0	0	0	0	I	0.0	0.0	0:00.00	rcu_bh
10	root	rt	0	0	0	0	S	0.0	0.0	0:00.10	migration/0
11	root	rt	0	0	0	0	S	0.0	0.0	0:00.84	watchdog/0
12	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cpuhp/0
13	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cpuhp/1
14	root	rt	0	0	0	0	S	0.0	0.0	0:01.05	watchdog/1
15	root	rt	0	0	0	0	S	0.0	0.0	0:00.09	migration/1
16	root	20	0	0	0	0	S	0.0	0.0	0:00.88	ksoftirqd/1
18	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/1:+

Exercise-9

Use cgroups to limit the cpu usage of a process

NOTE: *Perform below exercise on dnode1*

- Create 2 cgroups cpulimited and lesscpulimited

```
cgcreate -g cpu:/cpulimited  
cgcreate -g cpu:/lesscpulimited
```

```
root@dnode1:~# cgcreate -g cpu:/cpulimited  
[1]+  Done                  cpulimit -l 50 primes 0 9999999999 > /dev/null  
root@dnode1:~# cgcreate -g cpu:/lesscpulimited  
root@dnode1:~#
```

- Limit the cpu usage to 10% for cpulimited

```
cgset -r cpu.cfs_period_us=1000000 cpulimited/  
cgset -r cpu.cfs_quota_us=100000 cpulimited/
```

```
root@dnode1:~# cgset -r cpu.cfs_period_us=1000000 cpulimited/  
root@dnode1:~# cgset -r cpu.cfs_quota_us=100000 cpulimited/  
root@dnode1:~#
```

- Limit the cpu usage to 40% for lesscpulimited group

```
cgset -r cpu.cfs_period_us=1000000 lesscpulimited/  
cgset -r cpu.cfs_quota_us=400000 lesscpulimited/
```

```
root@dnode1:~# cgset -r cpu.cfs_period_us=1000000 lesscpulimited/  
root@dnode1:~# cgset -r cpu.cfs_quota_us=400000 lesscpulimited/  
root@dnode1:~#
```

- Run the primes process with these new cgroups

```
cgexec -g cpu:lesscpulimited primes 0 9999999999 > /dev/null &  
cgexec -g cpu:cpulimited primes 0 9999999999 > /dev/null &
```

```
root@dnode1:~# cgexec -g cpu:lesscpulimited primes 0 999999999 > /dev/null &
[1] 23063
root@dnode1:~# cgexec -g cpu:cpulimited primes 0 999999999 > /dev/null &
[2] 23064
root@dnode1:~#
```

- From VNC, observe from top command that 2 processes are limited to 40% and 10% of CPU respectively

desktop-kube-test (root's X desktop (ubuntu:1)) - VNC Viewer

Applications Home root@dnode1: /home/ubuntu 06:04 Ubuntu

root@dnode1: /home/ubuntu

File Edit View Search Terminal Help

top - 06:04:02 up 1 day, 1:02, 2 users, load average: 2.81, 1.04, 0.42
 Tasks: 102 total, 5 running, 56 sleeping, 0 stopped, 0 zombie
 %Cpu(s): 50.7 us, 0.7 sy, 42.7 ni, 4.5 id, 0.0 wa, 0.0 hi, 0.0 si, 1.5 st
 KiB Mem : 2043264 total, 1210012 free, 190656 used, 642596 buff/cache
 KiB Swap: 0 total, 0 free, 0 used, 1619808 avail Mem

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
23126	root	39	19	130732	70024	42208	R	88.0	3.4	0:17.37	unattended+
22762	root	20	0	61560	38864	5000	R	53.0	1.9	1:45.49	perl
23063	root	20	0	61124	38872	5012	R	40.0	1.9	0:28.24	perl
23064	root	20	0	61560	38856	4996	R	10.0	1.9	0:06.84	perl
22763	root	9	-11	8612	1132	1012	S	0.3	0.1	0:00.35	cpulimit
1	root	20	0	37876	5828	3916	S	0.0	0.3	0:18.05	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.01	kthreadd
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/0:+
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	mm_percpu+
7	root	20	0	0	0	0	S	0.0	0.0	0:00.48	ksoftirqd/0
8	root	20	0	0	0	0	I	0.0	0.0	0:04.85	rcu_sched
9	root	20	0	0	0	0	I	0.0	0.0	0:00.00	rcu_bh
10	root	rt	0	0	0	0	S	0.0	0.0	0:00.10	migration/0
11	root	rt	0	0	0	0	S	0.0	0.0	0:00.84	watchdog/0
12	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cpuhp/0
13	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cpuhp/1
14	root	rt	0	0	0	0	S	0.0	0.0	0:01.05	watchdog/1

Exercise-10

Explore the cgroup filesystem

NOTE: *Perform the below exercise on dnode1*

- All the cgroup controllers will be under /sys/fs/cgroup/

```
cd /sys/fs/cgroup/
ls
```

```
root@dnode1:~# cd /sys/fs/cgroup/
root@dnode1:/sys/fs/cgroup# ls
blkio cpu cpuacct cpu,cpuacct cpuset devices freezer hugetlb memory net_cls net_cls,net_prio net_prio perf_event pids rdma systemd
```

- Verify the default cpu, memory values for bash process

```
ps $$
```

```
root@dnode1:/sys/fs/cgroup# ps $$
  PID TTY          STAT       TIME COMMAND
22513 pts/1        Ss          0:00 -bash
root@dnode1:/sys/fs/cgroup#
```

```
cat /proc/<PID>/cgroup
```

```
root@dnode1:/sys/fs/cgroup# cat /proc/22513/cgroup
12:cpu,cpuacct:/user.slice/user-0.slice/session-1149.scope
11:hugetlb:/
10:freezer:/
9:perf_event:/
8:blkio:/user.slice/user-0.slice/session-1149.scope
7:rdma:/
6:memory:/user.slice/user-0.slice/session-1149.scope
5:net_cls,net_prio:/
4:devices:/user.slice/user-0.slice/session-1149.scope
3:cpuset:/
2:pids:/user.slice/user-0.slice/session-1149.scope
1:name=systemd:/user.slice/user-0.slice/session-1149.scope
root@dnode1:/sys/fs/cgroup#
```

(Verify the cgroup hierarchy followed for this process)

- Get the process ID from above 2 prime process (re-run if closed)
- Verify their cgroup file under /proc/PID/cgroup using their actual process IDs

Docker Part-I

Exercise-11

Docker Installation

NOTE: *Perform the below exercise on dnode1*

Setup the repository

- Update the apt package index:

```
apt-get update
```

```
Ubuntu 16.04.4 LTS
Welcome to Ubuntu 16.04.4 LTS (GNU/Linux 4.14.0-041400-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage

Get cloud support with Ubuntu Advantage Cloud Guest:
http://www.ubuntu.com/business/services/cloud

54 packages can be updated.
1 update is a security update.

Last login: Fri Jul 20 09:34:00 2018 from 192.168.122.1
root@dnode1:~# apt-get update
Hit:1 http://archive.ubuntu.com/ubuntu xenial InRelease
Get:2 http://security.ubuntu.com/ubuntu xenial-security InRelease [107 kB]
Get:3 http://archive.ubuntu.com/ubuntu xenial-updates InRelease [109 kB]
Get:4 http://archive.ubuntu.com/ubuntu xenial-backports InRelease [107 kB]
Fetched 323 kB in 0s (394 kB/s)
Reading package lists... Done
root@dnode1:~#
```

- Install packages to allow apt to use a repository over HTTPS:

```
apt-get install apt-transport-https ca-certificates curl software-properties-  
common
```

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
ca-certificates is already the newest version (20170717~16.04.1).
curl is already the newest version (7.47.0-1ubuntu2.8).
software-properties-common is already the newest version (0.96.20.7).
The following packages will be upgraded:
  apt-transport-https
1 upgraded, 0 newly installed, 0 to remove and 52 not upgraded.
Need to get 26.1 kB of archives.
After this operation, 0 B of additional disk space will be used.
Do you want to continue? [Y/n]
```

- Add Dockers official GPG key:

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
```

```
root@dnode1:~# curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
OK
root@dnode1:~#
```

- Verify that you now have the key with the fingerprint 9DC8 5822 9FC7 DD38 854A E2D8 8D81 803C 0EBF CD88, by searching for the last 8 characters of the fingerprint.

```
apt-key fingerprint 0EBFCD88
```

```
root@dnode1:~# apt-key fingerprint 0EBFCD88
pub 4096R/0EBFCD88 2017-02-22
    Key fingerprint = 9DC8 5822 9FC7 DD38 854A E2D8 8D81 803C 0EBF CD88
uid          Docker Release (CE deb) <docker@docker.com>
sub 4096R/F273FCD8 2017-02-22

root@dnode1:~#
```

- Use the following command to set up the stable repository.

```
add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
```

```
root@dnode1:~# add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
root@dnode1:~#
```


Install Docker CE

- Update the apt package index.

```
apt-get update
```

```
root@dnode1:~# apt-get update
Get:1 http://security.ubuntu.com/ubuntu xenial-security InRelease [107 kB]
Hit:2 http://archive.ubuntu.com/ubuntu xenial InRelease
Get:3 http://archive.ubuntu.com/ubuntu xenial-updates InRelease [109 kB]
Get:4 https://download.docker.com/linux/ubuntu xenial InRelease [65.8 kB]
Get:5 http://archive.ubuntu.com/ubuntu xenial-backports InRelease [107 kB]
Get:6 https://download.docker.com/linux/ubuntu xenial/stable amd64 Packages [3,904 B]
Fetched 393 kB in 0s (450 kB/s)
Reading package lists... Done
root@dnode1:~#
```

- Install the latest version of Docker CE

```
apt-get install docker-ce
```

```
root@dnode1:~# docker --version
Docker version 18.06.0-ce, build 0ffa825
root@dnode1:~#
```

- Verify docker version

```
docker --version
```

```
root@dnode1:~# docker --version
Docker version 18.06.0-ce, build 0ffa825
root@dnode1:~#
```

- Verify that Docker CE is installed correctly by running the hello-world image.

```
docker run hello-world
```

```

root@dnode1:~# docker run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
9db2ca6ccae0: Pull complete
Digest: sha256:4b8ff392a12ed9ea17784bd3c9a8b1fa3299cac44aca35a85c90c5e3c7afacdc
Status: Downloaded newer image for hello-world:latest

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
 1. The Docker client contacted the Docker daemon.
 2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
    (amd64)
 3. The Docker daemon created a new container from that image which runs the
    executable that produces the output you are currently reading.
 4. The Docker daemon streamed that output to the Docker client, which sent it
    to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/engine/userguide/

root@dnode1:~#

```

- List all the docker containers

```
docker ps -a
```

- List all the docker images

```
docker image ls
```

```

root@dnode1:~# docker image ls
REPOSITORY          TAG             IMAGE ID        CREATED         SIZE
hello-world         latest         2cb0d9787c4d   11 days ago    1.85kB

```

- Hello-world program is a basic scratch image. Check the source code

<https://github.com/docker-library/hello-world/blob/master/hello.c>

- Docker Inspect returns low level information on docker objects

```
docker inspect <container ID>
```

Exercise-12

Bring up Ubuntu Docker Container

NOTE: *Perform the below exercise on dnode1*

- Bring up Ubuntu container

```
docker run -it ubuntu bash
```

```
root@dnode1:~# docker run -it ubuntu bash
Unable to find image 'ubuntu:latest' locally
latest: Pulling from library/ubuntu
7996ebd2246a: Pull complete
de532f9a4f9f: Pull complete
7de2709b2a83: Pull complete
70b6ac64a142: Pull complete
23caf550e032: Pull complete
Digest: sha256:30e04ddada6eb09c12330c7df72cad1573916c7100168c34076808169ff6d805
Status: Downloaded newer image for ubuntu:latest
root@c604b837967e:/#
```

- List the directories

```
ls
```

```
root@c604b837967e:/# ls
bin boot dev etc home lib lib64 media mnt opt proc root run sbin srv sys usr var
root@c604b837967e:/#
```

- List the processes running

```
ps -ef
```

```
root@c604b837967e:/# ps -ef
UID          PID    PPID  C STIME TTY          TIME CMD
root           1        0  0  10:30 pts/0        00:00:00 bash
root          13         1  0  10:34 pts/0        00:00:00 ps -ef
root@c604b837967e:/#
```

- Check the Interfaces

```
ifconfig
```

(You will see it says command is not available.)

```
root@c604b837967e:/# ifconfig
bash: ifconfig: command not found
root@c604b837967e:/#
```

- Install relevant packages

```
apt-get update

apt-get install net-tools
```

```
root@c604b837967e:/# apt-get update
Get:1 http://archive.ubuntu.com/ubuntu bionic InRelease [242 kB]
Get:2 http://security.ubuntu.com/ubuntu bionic-security InRelease [83.2 kB]
Get:3 http://security.ubuntu.com/ubuntu bionic-security/universe Sources [9268 B]
Get:4 http://archive.ubuntu.com/ubuntu bionic-updates InRelease [88.7 kB]
Get:5 http://security.ubuntu.com/ubuntu bionic-security/multiverse amd64 Packages [1074 B]
Get:6 http://security.ubuntu.com/ubuntu bionic-security/universe amd64 Packages [44.4 kB]
Get:7 http://archive.ubuntu.com/ubuntu bionic-backports InRelease [74.6 kB]
Get:8 http://security.ubuntu.com/ubuntu bionic-security/main amd64 Packages [154 kB]
Get:9 http://archive.ubuntu.com/ubuntu bionic/universe Sources [11.5 MB]
Get:10 http://archive.ubuntu.com/ubuntu bionic/restricted amd64 Packages [13.5 kB]
Get:11 http://archive.ubuntu.com/ubuntu bionic/universe amd64 Packages [11.3 MB]
Get:12 http://archive.ubuntu.com/ubuntu bionic/main amd64 Packages [1344 kB]
Get:13 http://archive.ubuntu.com/ubuntu bionic/multiverse amd64 Packages [186 kB]
Get:14 http://archive.ubuntu.com/ubuntu bionic-updates/universe Sources [51.8 kB]
Get:15 http://archive.ubuntu.com/ubuntu bionic-updates/multiverse amd64 Packages [3679 B]
Get:16 http://archive.ubuntu.com/ubuntu bionic-updates/main amd64 Packages [322 kB]
Get:17 http://archive.ubuntu.com/ubuntu bionic-updates/universe amd64 Packages [163 kB]
Get:18 http://archive.ubuntu.com/ubuntu bionic-backports/universe amd64 Packages [2807 B]
Fetched 25.6 MB in 4s (7120 kB/s)
Reading package lists... Done
root@c604b837967e:/# apt-get install net-tools
```

- Check the interfaces again

```
ifconfig
```

```
root@d36986b4e237:/# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 172.17.0.2  netmask 255.255.0.0  broadcast 172.17.255.255
    ether 02:42:ac:11:00:02  txqueuelen 0  (Ethernet)
    RX packets 1569  bytes 25948682 (25.9 MB)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 1255  bytes 89446 (89.4 KB)
    TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING>  mtu 65536
    inet 127.0.0.1  netmask 255.0.0.0
    loop  txqueuelen 1000  (Local Loopback)
    RX packets 0  bytes 0 (0.0 B)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 0  bytes 0 (0.0 B)
    TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0

root@d36986b4e237:/# root@dnodel:~#
```

Criterion Networks

- SSH into dnode1, i.e, duplicate the browser or get a new dnode1 terminal from access devices. To get container ID use command "docker ps" in dnode1

```
docker ps
```

```
root@dnode1:~# docker ps
CONTAINER ID   IMAGE     COMMAND                  CREATED        STATUS        PORTS   NAMES
d36986b4e237   ubuntu   "bash"                   4 minutes ago  Up 4 minutes                unruffled_kilby
```

- Verify the IP address of container from host

```
docker inspect --format '{{ .NetworkSettings.IPAddress }}' <CONTAINERID>
```

```
root@dnode1:~# docker inspect --format '{{ .NetworkSettings.IPAddress }}' d36986b4e237
172.17.0.2
root@dnode1:~#
```

- To keep the docker running but detach from it gracefully, press CTRL+p, CTRL+q

Note: In case you exit the container, all the changes done to it, will be gone.

Exercise-13

Verify the PID, NET namespaces and CGROUP of ubuntu-bash process

- Get the PID of the bash process inside Ubuntu container

```
docker inspect --format '{{ .State.Pid }}' <CONTAINERID>
```

```
root@dnode1:~# docker inspect --format '{{ .State.Pid }}' d36986b4e237
24983
root@dnode1:~#
```

(say the output process number is [PIDbash])

- Verify the PID and NET namespace of INIT process on base node

```
readlink /proc/1/ns/pid
```

```
readlink /proc/1/ns/net
```

```
root@dnode1:~# readlink /proc/1/ns/pid
pid:[4026531836]
root@dnode1:~# readlink /proc/1/ns/net
net:[4026531993]
root@dnode1:~#
```

- Verify the PID and NET namespace of containers bash process

```
readlink /proc/<PIDbash>/ns/pid
```

```
readlink /proc/<PIDbash>/ns/net
```

```
root@dnode1:~# readlink /proc/24983/ns/pid
pid:[4026532205]
root@dnode1:~# readlink /proc/24983/ns/net
net:[4026532207]
root@dnode1:~#
```

- Verify the CPUINFO of base node

```
cat /proc/cpuinfo
```

Criterion Networks

```
processor      : 1
vendor_id     : GenuineIntel
cpu family    : 6
model         : 6
model name    : QEMU Virtual CPU version 2.2.0
stepping      : 3
microcode     : 0x1
cpu MHz       : 2299.808
cache size    : 4096 KB
physical id   : 1
siblings      : 1
core id       : 0
cpu cores     : 1
apicid        : 1
initial apicid : 1
fpu           : yes
fpu_exception : yes
cpuid level   : 4
wp            : yes
flags         : fpu de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pse36 clflush mmx fxsr sse sse2 syscall nx lm rep_good nopl cpuid pni vmx
               cx16 x2apic popcnt hypervisor lahf_lm abm tpr_shadow flexpriority ept
bugs          :
bogomips      : 4599.61
clflush size  : 64
cache_alignment : 64
address sizes  : 40 bits physical, 48 bits virtual
power management:

root@dnodel1:~#
```

- Verify the MEMINFO of base node

```
cat /proc/meminfo
```

```
Unevictable:      3652 kB
Mlocked:          3652 kB
SwapTotal:         0 kB
SwapFree:          0 kB
Dirty:             0 kB
Writeback:         0 kB
AnonPages:         81460 kB
Mapped:            97984 kB
Shmem:             21176 kB
Slab:              141436 kB
SReclaimable:      97688 kB
SUnreclaim:        43748 kB
KernelStack:       2400 kB
PageTables:        2976 kB
NFS_Unstable:       0 kB
Bounce:            0 kB
WritebackTmp:       0 kB
CommitLimit:      1021632 kB
Committed_AS:      523000 kB
VmallocTotal:     34359738367 kB
VmallocUsed:        0 kB
VmallocChunk:       0 kB
HardwareCorrupted:  0 kB
AnonHugePages:      0 kB
ShmemHugePages:     0 kB
ShmemPmdMapped:     0 kB
CmaTotal:           0 kB
CmaFree:            0 kB
HugePages_Total:    0
HugePages_Free:     0
HugePages_Rsvd:     0
HugePages_Surp:     0
Hugepagesize:       2048 kB
DirectMap4k:        83840 kB
DirectMap2M:        2013184 kB
root@dnode1:~#
```

- Verify the CPUSSET and MEMORY Hierarchy Path for this process

```
cat /proc/<PIDbash>/cgroup
```



```
root@dnode1:~# cat /proc/24983/cgroup
12:cpu,cpuacct:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
11:hugetlb:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
10:freezer:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
9:perf_event:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
8:blkio:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
7:rdma:/
6:memory:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
5:net_cls,net_prio:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
4:devices:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
3:cpuset:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
2:pids:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
1:name=systemd:/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334
root@dnode1:~#
```

(Check the cpuset row for CGROUP Hierarchy location)

- For CPUSET

```
cat /sys/fs/cgroup/cpuset/<PATH-TO-CGROUP>/cpuset.cpus
```

```
root@dnode1:~# cat /sys/fs/cgroup/cpuset/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334/cpuset.cpus
0-1
root@dnode1:~#
```

- For memory

```
cat /sys/fs/cgroup/memory/<PATH-TO-CGROUP>/memory.max_usage_in_bytes
```

```
root@dnode1:~# cat /sys/fs/cgroup/memory/docker/d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334/memory.max_usage_in_bytes
102592512
root@dnode1:~#
```

Exercise-14

Docker Networking Single Node with user defined network

- Perform below commands on dnode1
- Verify the default docker networks

```
docker network ls
```

```
root@dnode1:~# docker network ls
NETWORK ID          NAME                DRIVER              SCOPE
4d72e165c8be        bridge             bridge              local
85f2fffb1e6f        host               host                local
ca4e99008c59        none               null                local
root@dnode1:~#
```

- Inspect the docker default network bridge

```
docker network inspect bridge
```

```
[
  {
    "Name": "bridge",
    "Id": "4d72e165c8befaee96a6f2b64b85e4b6ca30873e14f0d0a6c1ecbae425cf9e5d",
    "Created": "2018-07-22T10:20:21.421773616Z",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "172.17.0.0/16"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
    "Containers": {
      "d36986b4e23760e22263da8de80a76fbb5d38bca66afb2e2f0b2279f14534334": {
        "Name": "unruffled_kilby",
        "EndpointID": "44066613c7b89cbf23c6a716cd029d6c12f93e2972034c1697450aae886e5ef1",
        "MacAddress": "02:42:ac:11:00:02",
        "IPv4Address": "172.17.0.2/16",
        "IPv6Address": ""
      }
    },
    "Options": {
      "com.docker.network.bridge.default_bridge": "true",

```

Criterion Networks

- Create your own bridge network

```
docker network create -d bridge mynet
```

```
root@dnode1:~# docker network create -d bridge mynet
5ec2bd9e688bdc8e88c6f126932bf5763b871a21f163e4c3e469a1f9e810c946
```

- List the networks and inspect new network

```
docker network ls
```

```
root@dnode1:~# docker network ls
NETWORK ID          NAME                DRIVER              SCOPE
4d72e165c8be        bridge             bridge              local
85f2ffffb1e6f       host               host                local
5ec2bd9e688b        mynet              bridge              local
ca4e99008c59        none               null                local
root@dnode1:~#
```

```
docker network inspect mynet
```

```

root@dnodel1:~# docker network inspect mynet
[
  {
    "Name": "mynet",
    "Id": "5ec2bd9e688bdc8e88c6f126932bf5763b871a21f163e4c3e469a1f9e810c94",
    "Created": "2018-07-22T11:02:31.739347018z",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": {},
      "Config": [
        {
          "Subnet": "172.18.0.0/16",
          "Gateway": "172.18.0.1"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
    "Containers": {},
    "Options": {},
    "Labels": {}
  }
]
root@dnodel1:~#

```

- Create 2 new docker containers with busybox (Linux image with most of network tools already installed) image

```
docker run -itd --net=mynet --name box1 busybox
```

```
docker run -itd --net=mynet --name box2 busybox
```

```

root@dnodel1:~# docker run -itd --net=mynet --name box1 busybox
Unable to find image 'busybox:latest' locally
latest: Pulling from library/busybox
75a0e65efd51: Pull complete
Digest: sha256:d21b79794850b4b15d8d332b451d95351d14c951542942a816eea69c9e04b240
Status: Downloaded newer image for busybox:latest
f72ef52a58cca13451ab5a51dfbc13e10d837b5a545f2c4dfdb67e84e2d3dc88
root@dnodel1:~# docker run -itd --net=mynet --name box2 busybox
bf6ee9ceed700fb1250e3000678e9487488d282b1f56aacb835d383488a08db1
root@dnodel1:~#

```

- Attach to box1 container and observe its IP address (say box1-IP)

```
docker attach box1
```

```
ifconfig
```

```
root@dnodel:~# docker attach box1
/ # ifconfig
eth0      Link encap:Ethernet  HWaddr 02:42:AC:12:00:02
          inet addr:172.18.0.2  Bcast:172.18.255.255  Mask:255.255.0.0
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:18 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:1452 (1.4 KiB)  TX bytes:0 (0.0 B)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          UP LOOPBACK RUNNING  MTU:65536  Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)

/ #
```

Press Ctrl+p & Ctrl+q to detach from the box1 container

- Attach to box2 container and ping "box1" IP address. Verify it is reachable.

```
docker attach box2
```

```
ifconfig
```

```
ping <box1-IP>
```

```

root@dnode1:~# docker attach box2
/ # ifconfig
eth0      Link encap:Ethernet  HWaddr 02:42:AC:12:00:03
          inet addr:172.18.0.3  Bcast:172.18.255.255  Mask:255.255.0.0
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:14 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:1076 (1.0 KiB)  TX bytes:0 (0.0 B)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          UP LOOPBACK RUNNING  MTU:65536  Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)

/ # ping -c3 172.18.0.2
PING 172.18.0.2 (172.18.0.2): 56 data bytes
64 bytes from 172.18.0.2: seq=0 ttl=64 time=0.114 ms
64 bytes from 172.18.0.2: seq=1 ttl=64 time=0.089 ms
64 bytes from 172.18.0.2: seq=2 ttl=64 time=0.089 ms

--- 172.18.0.2 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max = 0.089/0.097/0.114 ms
/ #

```

Exercise-15

Explore Overlay2 storage driver

- Create 2 ubuntu containers
- Inspect details of 2 containers

```
docker inspect <CONTAINERID-1>
```

```
docker inspect <CONTAINERID-2>
```

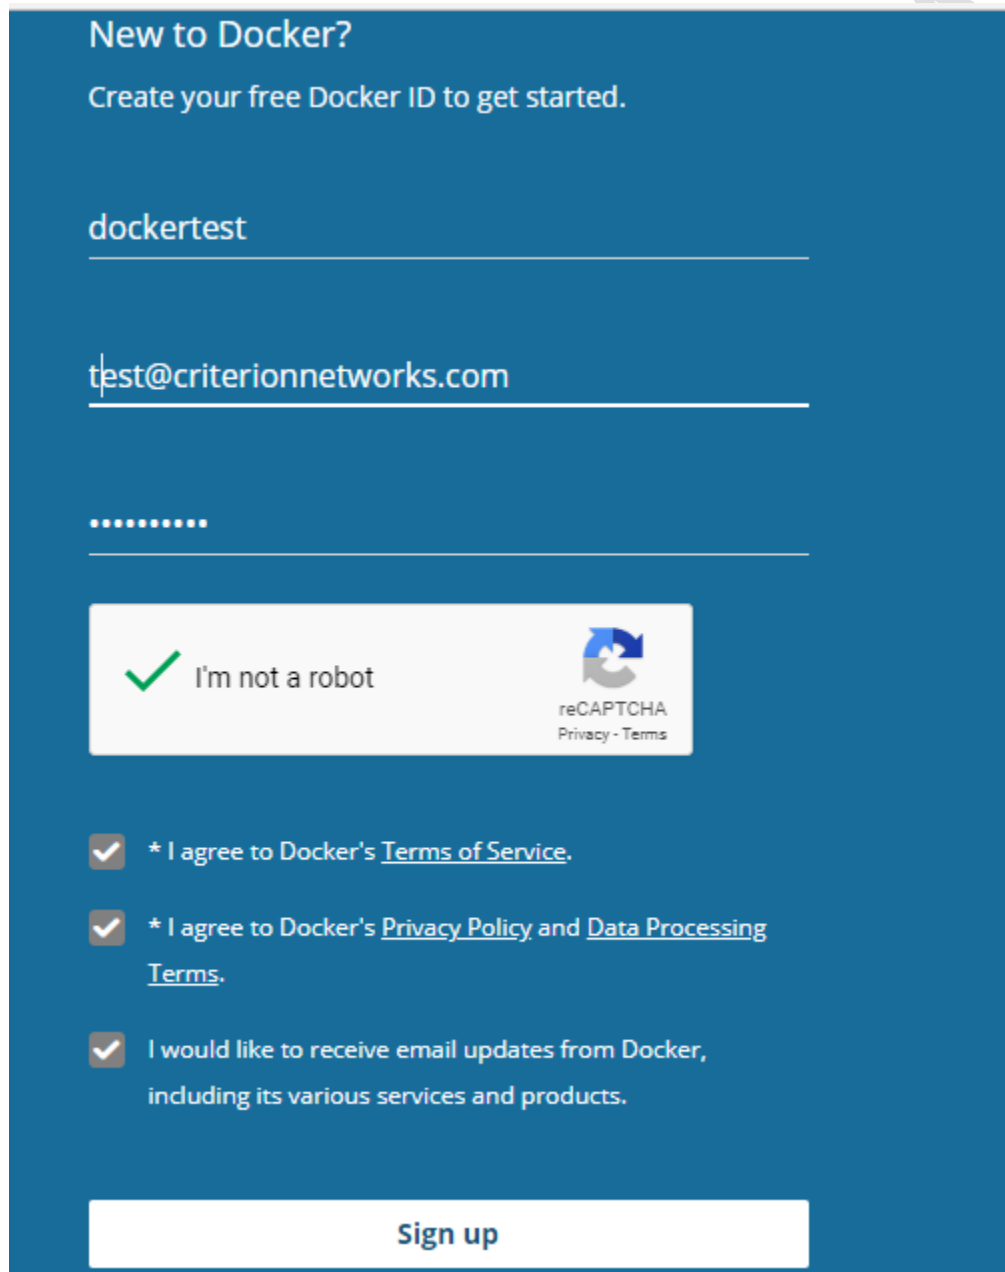
- Verify the Graph Driver details for both containers
- Observe the similarity in both the graph driver outputs
- Observe the directory structure for overlay2 driver i.e lower, diff, merged, work

Docker Part-II

Exercise-16

Sign-up for Docker ID

- Visit <https://cloud.docker.com>
- Sign up to get a Docker ID
- Verify you are able to login with Docker ID and password

A screenshot of the Docker sign-up page. The background is a solid blue color. At the top, the text "New to Docker?" is in white, followed by "Create your free Docker ID to get started." Below this, there are three input fields: the first contains "dockertest", the second contains "test@critterionnetworks.com", and the third contains a series of dots representing a password. Below the password field is a white box containing a green checkmark and the text "I'm not a robot", along with a reCAPTCHA logo and the text "reCAPTCHA Privacy - Terms". Below this box are three checkboxes, all of which are checked. The first checkbox is followed by the text "* I agree to Docker's Terms of Service.", the second by "* I agree to Docker's Privacy Policy and Data Processing Terms.", and the third by "I would like to receive email updates from Docker, including its various services and products." At the bottom of the form is a white button with the text "Sign up" in blue.

New to Docker?


Create your free Docker ID to get started.

dockertest

test@critterionnetworks.com

.....

☒ I'm not a robot

 reCAPTCHA
Privacy - Terms

☒ * I agree to Docker's [Terms of Service](#).

☒ * I agree to Docker's [Privacy Policy](#) and [Data Processing Terms](#).

☒ I would like to receive email updates from Docker, including its various services and products.

Sign up

Exercise-17

Build the Apache2 Web Server Container using DockerFile

- Create a new directory learning/apache2

```
mkdir -p learning/apache2
```

```
cd learning/apache2
```

```
root@dnodel1:~# mkdir -p learning/apache2
root@dnodel1:~# cd learning/apache2
root@dnodel1:~/learning/apache2#
```

- Create a new file named Dockerfile

```
vim Dockerfile
```

- Insert following lines in the file, save and close it

```
FROM ubuntu:14.04

RUN apt-get update

RUN apt-get install apache2 -y

EXPOSE 80
```

```
FROM ubuntu:latest
RUN apt-get update
RUN apt-get install apache2 -y
EXPOSE 80
~
~
~
~
~
~
```

- Build the docker image

```
docker build -t apache2 .
```

```
root@dnode1:~/learning/apache2# docker build -t apache2 .
Sending build context to Docker daemon 2.048kB
Step 1/4 : FROM ubuntu:latest
--> 74f8760a2a8b
Step 2/4 : RUN apt-get update
--> Running in f7c3c3c2b3e3
Get:1 http://security.ubuntu.com/ubuntu bionic-security InRelease [83.2 kB]
Get:2 http://archive.ubuntu.com/ubuntu bionic InRelease [242 kB]
Get:3 http://security.ubuntu.com/ubuntu bionic-security/universe Sources [9268 B]
Get:4 http://archive.ubuntu.com/ubuntu bionic-updates InRelease [88.7 kB]
Get:5 http://security.ubuntu.com/ubuntu bionic-security/main amd64 Packages [154 kB]
Get:6 http://archive.ubuntu.com/ubuntu bionic-backports InRelease [74.6 kB]
Get:7 http://security.ubuntu.com/ubuntu bionic-security/multiverse amd64 Packages [1074 B]
Get:8 http://security.ubuntu.com/ubuntu bionic-security/universe amd64 Packages [44.4 kB]
Get:9 http://archive.ubuntu.com/ubuntu bionic/universe Sources [11.5 MB]
Get:10 http://archive.ubuntu.com/ubuntu bionic/multiverse amd64 Packages [186 kB]
Get:11 http://archive.ubuntu.com/ubuntu bionic/restricted amd64 Packages [13.5 kB]
Get:12 http://archive.ubuntu.com/ubuntu bionic/main amd64 Packages [1344 kB]
Get:13 http://archive.ubuntu.com/ubuntu bionic/universe amd64 Packages [11.3 MB]
Get:14 http://archive.ubuntu.com/ubuntu bionic-updates/universe Sources [51.8 kB]
Get:15 http://archive.ubuntu.com/ubuntu bionic-updates/multiverse amd64 Packages [3679 B]
Get:16 http://archive.ubuntu.com/ubuntu bionic-updates/main amd64 Packages [322 kB]
Get:17 http://archive.ubuntu.com/ubuntu bionic-updates/universe amd64 Packages [163 kB]
Get:18 http://archive.ubuntu.com/ubuntu bionic-backports/universe amd64 Packages [2807 B]
Fetched 25.6 MB in 4s (7148 kB/s)
Reading package lists...
```

- Verify the image is built

```
docker image ls
```

```
root@dnode1:~/learning/apache2# docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
apache2	latest	9c549532a4fe	4 seconds ago	219MB
ubuntu	latest	74f8760a2a8b	5 days ago	82.4MB
busybox	latest	22c2dd5ee85d	5 days ago	1.16MB
hello-world	latest	2cb0d9787c4d	11 days ago	1.85kB

```
root@dnode1:~/learning/apache2#
```

- Run the docker Image and verify the functionality

```
docker run -p 8000:80 -it apache2 bash
```

```
service apache2 restart
```

```
root@dnode1:~/learning/apache2# docker run -p 8000:80 -it apache2 bash
root@ab3edaca72e4:/# service apache2 restart
* Restarting Apache httpd web server apache2
AH00558: apache2: Could not reliably determine the server's fully qualified domain name, using 172.17.0.3. Set the 'ServerName' directive globally to suppress this message
root@ab3edaca72e4:/#
```

NOTE: Do not exit the container and open another dnode1 terminal to view the running apache2 container

- Verify the docker is running at proper port on host dnode1 (The below command has to be performed on the other donode1 terminal.)

```
docker ps
```

Criterion Networks

```
root@dnode1:~/learning/apache2# docker ps
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS                    NAMES
ab3edaca72e4   apache2   "bash"                  56 seconds ago Up 55 seconds  0.0.0.0:8000->80/tcp    upbeat_shannon
bf6ee9ceed70   busybox   "sh"                   21 minutes ago Up 21 minutes             box2
f72ef52a58cc   busybox   "sh"                   22 minutes ago Up 21 minutes             box1
d36986b4e237   ubuntu    "bash"                 About an hour ago Up About an hour          unruffled_kilby
root@dnode1:~/learning/apache2#
```

- Get the IP address of dnode1 for interface "ens3"
- Open VNC on desktop1 machine. Open Firefox web browser
- Perform http request in the web browser

```
http://<IP-ADDR-dnode1-ens3>:8000
```

(Verify Apache2 web page is displayed)

Exercise-18

Build an application using Python Run-time

- Create a new directory learning/app

```
mkdir -p learning/app  
  
cd learning/app
```

```
root@dnodel:~# mkdir -p learning/app  
root@dnodel:~# cd learning/app
```

- Create a new file named Dockerfile

```
vim Dockerfile
```

- Insert following lines in the file, save and close it

```
# Use an official Python runtime as a parent image  
FROM python:2.7-slim  
  
# Set the working directory to /app  
WORKDIR /app  
  
# Copy the current directory contents into the container at /app  
ADD . /app  
  
# Install any needed packages specified in requirements.txt  
RUN pip install --trusted-host pypi.python.org -r requirements.txt
```

```
# Make port 80 available to the world outside this container

EXPOSE 80

# Define environment variable

ENV NAME World

# Run app.py when the container launches

CMD ["python", "app.py"]
```

```
# Use an official Python runtime as a parent image
FROM python:2.7-slim

# Set the working directory to /app
WORKDIR /app

# Copy the current directory contents into the container at /app
ADD . /app

# Install any needed packages specified in requirements.txt
RUN pip install --trusted-host pypi.python.org -r requirements.txt

# Make port 80 available to the world outside this container
EXPOSE 80

# Define environment variable
ENV NAME World

# Run app.py when the container launches
CMD ["python", "app.py"]
~
~
~
~
```

- Create a file named requirements.txt

```
vim requirements.txt
```

```
Flask
```

Redis

```
flask
Redis
```

```
~
~
~
~
~
~
~
~
~
```

- Create a file named app.py

```
root@dnodel:~/learning/app# vim app.py
```

```
from flask import Flask

from redis import Redis, RedisError

import os

import socket

# Connect to Redis

redis = Redis(host="redis", db=0, socket_connect_timeout=2, socket_timeout=2)

app = Flask(__name__)

@app.route("/")

def hello():

    try:

        visits = redis.incr("counter")

    except RedisError:
```

```

        visits = "<i>cannot connect to Redis, counter disabled</i>"

    html = "<h3>Hello {name}!</h3>" \
          "<b>Hostname:</b> {hostname}<br/>" \
          "<b>Visits:</b> {visits}"

    return html.format(name=os.getenv("NAME", "world"), hostname=socket.gethostname(), visits=visits)

if __name__ == "__main__":

    app.run(host='0.0.0.0', port=80)

```

```

from flask import Flask
from redis import Redis, RedisError
import os
import socket

# Connect to Redis
redis = Redis(host="redis", db=0, socket_connect_timeout=2, socket_timeout=2)

app = Flask(__name__)

@app.route("/")
def hello():
    try:
        visits = redis.incr("counter")
    except RedisError:
        visits = "<i>cannot connect to Redis, counter disabled</i>"

    html = "<h3>Hello {name}!</h3>" \
          "<b>Hostname:</b> {hostname}<br/>" \
          "<b>Visits:</b> {visits}"
    return html.format(name=os.getenv("NAME", "world"), hostname=socket.gethostname(), visits=visits)

if __name__ == "__main__":
    app.run(host='0.0.0.0', port=80)

```

- Build the docker image

```
docker build -t flaskapp .
```

```
Removing intermediate container 0603705c82d0
---> 87363ce05fbd
Step 5/7 : EXPOSE 80
---> Running in 510cedaf75c0
Removing intermediate container 510cedaf75c0
---> bc77460508c7
Step 6/7 : ENV NAME World
---> Running in b005ceac5fc5
Removing intermediate container b005ceac5fc5
---> 900c7d97549a
Step 7/7 : CMD ["python", "app.py"]
---> Running in d3e0aad051f6
Removing intermediate container d3e0aad051f6
---> f45e2495fcb5
Successfully built f45e2495fcb5
Successfully tagged flaskapp:latest
root@dnodel1:~/learning/app#
```

- Verify the image is built

```
docker image ls
```

```
root@dnodel1:~/learning/app# docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
flaskapp	latest	f45e2495fcb5	38 seconds ago	132MB
apache2	latest	9c549532a4fe	18 minutes ago	219MB
python	2.7-slim	42967d04ddc5	5 days ago	120MB
ubuntu	latest	74f8760a2a8b	5 days ago	82.4MB
busybox	latest	22c2dd5ee85d	5 days ago	1.16MB
hello-world	latest	2cb0d9787c4d	11 days ago	1.85kB

```
root@dnodel1:~/learning/app#
```

- Run the docker Image in detached mode and verify the functionality

```
docker run -d -p 4000:80 -it flaskapp
```

```
root@dnodel1:~/learning/app# docker run -d -p 4000:80 -it flaskapp
c87c322c36a5ca28d0f2dcc5b50c253ca07f21d3d52f9a1911fc307f86cb645c
root@dnodel1:~/learning/app#
```

- Verify the docker is running at proper port on host dnodel

```
docker ps
```

```
root@dnodel1:~/learning/app# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
ab3edaca72e4	apache2	"bash"	15 minutes ago	Up 15 minutes	0.0.0.0:8000->80/tcp	upbeat_shannon
bf6ee9ced70	busybox	"sh"	36 minutes ago	Up 36 minutes		box2
f72cf52a58cc	busybox	"sh"	36 minutes ago	Up 36 minutes		box1
d36986b4e237	ubuntu	"bash"	About an hour ago	Up About an hour		unruffled_kilby

```
root@dnodel1:~/learning/app#
```

- Get the IP address of dnodel for interface "ens3"

Criterion Networks

- Open VNC on desktop1 machine. Open Firefox web browser
- Perform http request in the web browser

```
http://<IP-ADDR-dnode1>:4000
```

(Verify Flask web page and Redis Error is displayed)

CONFIDENTIAL

Exercise-19

Push the Docker Images into Docker Hub Registry

- Login into docker hub using Docker ID

```
docker login
```

(Provide your docker username and password to login)

- Tag the image to be pushed into Public Repository
- Format: `docker tag imagename username/repository:tag`

NOTE: Replace `nareshthukkani` with your docker username in command below

Example:

```
docker tag flaskapp nareshthukkani/learning:flaskapp
```

- Verify docker images. Displays new tag image with same Image ID
- Push the image into Docker Hub

NOTE: Replace `nareshthukkani` with your username in command below

```
docker push nareshthukkani/learning:flaskapp
```

```
root@dnode1:~# docker login
Authenticating with existing credentials...
WARNING! Your password will be stored unencrypted in /home/ubuntu/.docker/config.json.
Configure a credential helper to remove this warning. See
https://docs.docker.com/engine/reference/commandline/login/#credentials-store

Login Succeeded
root@dnode1:~# docker tag flaskapp ngupta786/learning:flaskapp
root@dnode1:~# docker push ngupta786/learning:flaskapp
The push refers to repository [docker.io/ngupta786/learning]
f8a62ebecc5e: Pushed
5c9a807d5891: Pushed
4c4fde5866ee: Pushed
e8f274959cd1: Pushed
586cb5f00e7c: Pushed
21b2e9972483: Pushed
cdb3f9544e4c: Pushed
flaskapp: digest: sha256:a94285bc602ba527747c8f46cdafaa1117cf8329921ff97300254bae7221960a size: 1788
root@dnode1:~#
```

- Login into docker cloud account. Click on Repositories to check the image pushed
- Now, please push `apache2` image as well in the above way

Exercise-20

Docker Service using docker-compose

- Perform the below commands on "dnode1" node. We will bring up single node dswarm cluster.
- Create a file named docker-compose.yml

```
vim docker-compose.yml
```

```
root@dnode1:~# vim docker-compose.yml
root@dnode1:~#
```

- Insert the below lines [replace the "nareshthukkani" with your docker username]

```
version: "3"

services:

  web:

    # replace username/repo:tag with your name and image details

    image: nareshthukkani/learning:flaskapp

    deploy:

      replicas: 3

      resources:

        limits:

          cpus: "0.1"

          memory: 50M

      restart_policy:

        condition: on-failure

    ports:

      - "5000:80"

    networks:
```

```
- webnet
```

```
networks:
```

```
webnet:
```

```
version: "3"
services:
  web:
    # replace username/repo:tag with your name and image details
    image: nareshthukkani/learning:flaskapp
    deploy:
      replicas: 3
      resources:
        limits:
          cpus: "0.1"
          memory: 50M
      restart_policy:
        condition: on-failure
    ports:
      - "5000:80"
    networks:
      - webnet
networks:
  webnet:
~
~
~
```

- Initialize the Docker Swarm

```
docker swarm init --advertise-addr 192.168.122.205
```

```
root@dnode1:~# docker swarm init --advertise-addr 192.168.122.205
Swarm initialized: current node (n2g36wwrobqy9omlgy8ujtpx7) is now a manager.

To add a worker to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-37u2yvbjiyy79xite1ntywikyla2mrh0bwtj4v8fwnkaz86hfz-arbsjtg1kivzckdprjlkxkrir 192.168.122.205:2377

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.

root@dnode1:~#
```

- Deploy the stack

```
docker stack deploy -c docker-compose.yml getstartedlab
```

```
root@dnode1:~# docker stack deploy -c docker-compose.yml getstartedlab
Creating network getstartedlab_webnet
Creating service getstartedlab_web
root@dnode1:~#
```

- Verify the docker service is launched

```
docker service ls
```

```
root@dnode1:~# docker service ls
ID                NAME                MODE                REPLICAS            IMAGE                PORTS
rd4770xxcxez      getstartedlab_web    replicated           3/3                 nareshthukkani/learning:flaskapp    *:5000->80/tcp
root@dnode1:~#
```

NOTE: Fetch the port number exposed in the displayed service which will be used with the advertised address used in the init command.

- Verify the new containers deployed

```
docker ps
```

```
root@dnode1:~# docker ps
CONTAINER ID        IMAGE                COMMAND              CREATED             STATUS              PORTS              NAMES
83c4f9380c6d       nareshthukkani/learning:flaskapp    "python app.py"     38 seconds ago     Up 32 seconds      80/tcp             getstartedl
ab_web.3.tlpe5sq7ihdiran4gjlymxwm7    nareshthukkani/learning:flaskapp    "python app.py"     38 seconds ago     Up 33 seconds      80/tcp             getstartedl
b9e6383476d2       nareshthukkani/learning:flaskapp    "python app.py"     38 seconds ago     Up 35 seconds      80/tcp             getstartedl
ab_web.1.o6d44vvd53j38dq6sb24k04gs    nareshthukkani/learning:flaskapp    "python app.py"     38 seconds ago     Up 40 minutes      0.0.0.0:8000->80/tcp    upbeat_shan
bc9c17c8c10d       apache2              "bash"              40 minutes ago     Up 40 minutes
ab3edaca72e4       busybox              "sh"                About an hour ago   Up About an hour
b26es9cead70       busybox              "sh"                About an hour ago   Up About an hour
f72ef52a50cc       ubuntu              "bash"              2 hours ago        Up 2 hours
d36986b4e237       ilby
root@dnode1:~#
```

- Verify from VNC, browse the

```
http://<ip-address-advertised>:<port-fetched-in-service-display-command>
```

- Web Browser repeatedly to this service gives different web page every time
- Change the replicas value in docker-compose.yml to 5 and deploy the stack again
- Verify 5 containers are deployed
- Take down the app

```
docker stack rm getstartedlab
```

```
root@dnode1:~# docker stack rm getstartedlab
Removing service getstartedlab_web
Removing network getstartedlab_webnet
root@dnode1:~#
```

- Take down the swarm

```
docker swarm leave --force
```

```
root@dnode1:~# docker swarm leave --force
Node left the swarm.
root@dnode1:~#
```

Exercise-21

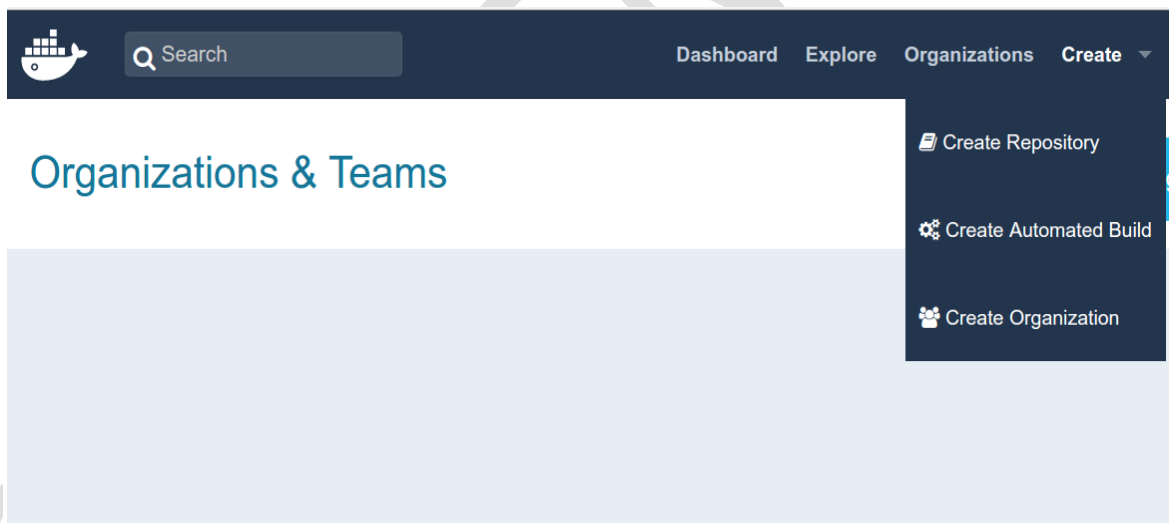
Docker with DevOps

1.Docker Github

- In this exercise we will use github , so you need to have a github account.You can use your existing github account or else create a new one.
- Sign-in to github and open following URL <https://github.com/rauhada/docker-training>
- Click on Fork , this will create your own copy in your account
- View all the files present in the keystone and Glance directories
- Now go to <http://hub.docker.com> and create an account

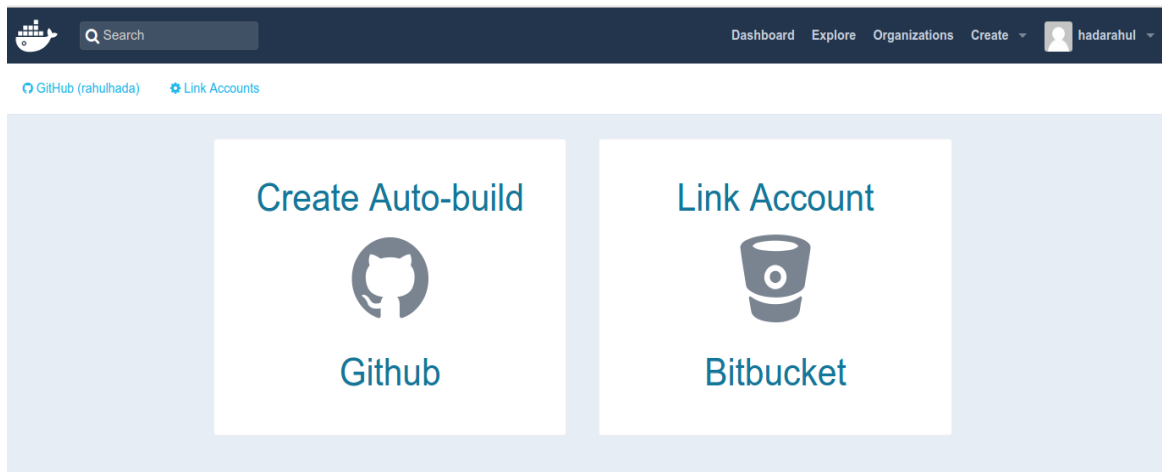
2.Build Docker Repo

- Now link the github docker-training repo with docker-hub
- Click on Create Automatic Build



3.Auto-Build for Git

- Create auto-build for github repo



4.Select Docker Training Repo

- Select docker-training repo and give name to your docker image. We can name it as docker-training-glance for glance image and docker-training-keystone for keystone image.
- Click on Create button



Create Automated Build

The screenshot shows the 'Create Automated Build' form. It has a light blue background. The form is divided into two main sections. The top section is titled 'Repository Namespace & Name*' and contains two input fields: one for the namespace (set to 'hadarahul') and one for the repository name (set to 'docker-training-glance'). To the right of these fields is a 'Visibility' dropdown menu set to 'public'. Below this section is a 'Short Description*' field with the text 'Glance Docker Image'. At the bottom of the form, there is a note: 'By default Automated Builds will match branch names to Docker build tags. [Click here to customize](#) behavior.' and a blue 'Create' button.

5.Build Setting

- After that click on Build Settings tab

Build Settings

☒ When active, builds will happen automatically on pushes.

The build rules below specify how to build your source into Docker images. The name can be a string or a regex. The Docker Tag name may contain variables. We currently support {sourcerefs}, which refers to the source branch/tag name. [Show more](#)

Source Repository: [rauhada/docker-training](#)

Type	Name	Dockerfile Location	Docker Tag Name	
Branch	master	/	latest	+
Branch	All branches except master	/	Same as branch	-

Trigger

- Here we need to give the location of Dockerfile directory. The Dockerfile located in glance and keystone directories.

Build Settings

☒ When active, builds will happen automatically on pushes.

The build rules below specify how to build your source into Docker images. The name can be a string or a regex. The Docker Tag name may contain variables. We currently support {sourcerefs}, which refers to the source branch/tag name. [Show more](#)

Source Repository: [rauhada/docker-training](#)



Type	Name	Dockerfile Location	Docker Tag Name	
Branch	master	/glance/	latest	+
Branch	All branches except master	/	Same as branch	-

Trigger

Save Changes

6. Verify Build Status

- Now click on **Save Changes** and then **trigger** to start building. Click on build details tab to see the build status.

Repo Info	Tags	Dockerfile	Build Details	Build Settings	Collaborators	Webhooks	Settings
Status	Actions	Tag	Created	Last Updated			
 Queued	 Cancel	latest	a few seconds ago	a few seconds ago			

- This will take few minutes to build and wait till the status of build is successful. Repeat the steps from 2 to 6 to build keystone image.

Exercise-22

To Spawn the Mysql Database, Keystone and Glance Container

On dswarm1 node

- Create an overlay network to connect hosted container together

```
docker network create \  
  
--driver overlay \  
  
--subnet 10.0.9.0/24 \  
  
My-multi-host-network
```

```
root@dswarm1:/root/learning# docker network create \  
> --driver overlay \  
> --subnet 10.0.9.0/24 \  
> My-multi-host-network  
14h0cwki9hpf1awhj2omtss0p  
root@dswarm1:/root/learning#
```

- To verify the created network

```
docker network ls
```

```
root@dswarm1:/root/learning# docker network ls  
NETWORK ID          NAME                                DRIVER                SCOPE  
14h0cwki9hpf        My-multi-host-network              overlay                swarm  
e071c6e92455        bridge                             bridge                 local  
2cbe9c70f5ca        docker_gwbridge                    bridge                 local  
fxbsfdbj84p3        getstartedlab_webnet               overlay                swarm  
772fe9803dba        host                                host                   local  
mirs077ko3ny        ingress                             overlay                swarm  
7ff0e75b2532        none                                null                   local  
root@dswarm1:/root/learning#
```

- Create a docker service for mysql

```
docker service create -e MYSQL_ROOT_PASSWORD=admin123 --replicas 1 --network  
getstartedlab_webnet --name mysql mysql:5.7.23
```

Criterion Networks

```
root@dswarm1:/root/learning# docker service create -e MYSQL_ROOT_PASSWORD=admin123 --replicas 1 --network My-multi-host-network --name mysql mysql
wksamzf94gt19cc0569oo4cpw
overall progress: 1 out of 1 tasks
1/1: running [=====>]
verify: Service converged
root@dswarm1:/root/learning#
```

- Verify mysql service is running and observe the node it is spawned

```
docker service ps mysql
```

```
root@dswarm1:/root/learning# docker service ps mysql
ID            NAME      IMAGE          NODE      DESIRED STATE   CURRENT STATE      ERROR      PORTS
m3o30n4vb96s  mysql.1   mysql:latest   dnode2    Running         Running 33 seconds ago
```

- Create a docker service for keystone (replace "nareshthukkani" with your docker username)

```
docker service create --replicas=1 --network getstartedlab_webnet --name keystone nareshthukkani/docker-training-keystone
```

```
root@dswarm1:/root/learning# docker service create --replicas=1 --network getstartedlab_webnet --name keystone nareshthukkani/docker-training-keystone
cbx510vsirl3x7tqetmy2em75
overall progress: 1 out of 1 tasks
1/1: running [=====>]
verify: Service converged
root@dswarm1:/root/learning#
```

- Verify keystone service is running and observe node it is spawned

```
docker service ps keystone
```

```
root@dswarm1:/root/learning# docker service ps keystone
ID            NAME      IMAGE          NODE      DESIRED STATE   CURRENT STATE      ERROR      PORTS
wp8zs08xjd0i  keystone.1  nareshthukkani/docker-training-keystone:latest  dnode1    Running         Running 43 seconds ago
```

- Create a docker service for glance (replace "nareshthukkani" with your docker username)

```
docker service create --replicas=1 --network getstartedlab_webnet --name glance nareshthukkani/docker-training-glance
```

```
root@dswarm1:/root/learning# docker service create --replicas=1 --network getstartedlab_webnet --name glance nareshthukkani/docker-training-glance
jmlbzdj06k36ybn6koehrq2p7
overall progress: 1 out of 1 tasks
1/1: running [=====>]
verify: Service converged
root@dswarm1:/root/learning#
```

- Verify glance service is running and observe node it is spawned

```
docker service ps glance
```

```
root@dswarm1:/root/learning# docker service ps glance
ID            NAME      IMAGE          NODE      DESIRED STATE   CURRENT STATE      ERROR      PORTS
traezfo2g3ec  glance.1   nareshthukkani/docker-training-glance:latest  dnode1    Running         Running about a minute ago
```

Exercise-23

To configure services and verify functionality

- Fetch the container ID of keystone container and login to the container

```
docker ps
```

```
docker exec -it <KEYSTONE-CONTAINERID> /bin/bash
```

```
root@dnode1:~# docker ps
CONTAINER ID        IMAGE
NAMES
1cd901da7a69       nareshthukkani/docker-training-glance:latest
glance.1.traezfo2g3ecy73fwvixyd9aq
ca0fb3dff0e0       nareshthukkani/docker-training-keystone:latest
p    keystone.1.wp8zs08xjd0if5enrlzgu3bv5
18dd0e282742       nareshthukkani/learning:flaskapp
getstartedlab_web.5.nbxcaynoph8qjkothhbwfupe9
cf6507b86c5f       nareshthukkani/learning:flaskapp
getstartedlab_web.3.wmew72a8ce5r5rrsqppasuakj
ab3edaca72e4       apache2
cp    upbeat_shannon
bf6ee9ceed70       busybox
box2
f72ef52a58cc       busybox
box1
d36986b4e237       ubuntu
unruffled_kilby
root@dnode1:~# docker exec -it ca0fb3dff0e0 /bin/bash
root@ca0fb3dff0e0:/#
```

- Create openrc file, paste the contents and save the file

```
vi openrc
```

```
root@ca0fb3dff0e0:/# vi openrc
root@ca0fb3dff0e0:/#
```

- Insert the following lines

```
export OS_TOKEN=token123
```

```
export OS_URL=http://keystone:35357/v2.0
```

```
export OS_TOKEN=token123
export OS_URL=http://keystone:35357/v2.0
~
~
~
~
~
~
~
~
```

- Source openrc file

```
source openrc
```

- Create an entry for identity service

```
openstack service create --name keystone --description "OpenStack Identity" identity
```

- Create Identity service endpoints

```
openstack endpoint create \
--publicurl http://keystone:5000/v2.0 \
--internalurl http://keystone:5000/v2.0 \
--adminurl http://keystone:35357/v2.0 \
--region RegionOne \
identity
```

- Create admin project

```
openstack project create admin
```

- Create admin user

```
openstack user create --password admin admin
```

- Create role admin

```
openstack role create admin
```

Criterion Networks

- Assign role admin to user admin in admin project

```
openstack role add --project admin --user admin admin
```

- Create project with name service

```
openstack project create --description "Service Project" service
```

- Create user with name glance

```
openstack user create --password GLANCE_PASS glance
```

- Assign admin role to glance user in service project

```
openstack role add --project service --user glance admin
```

- Create service with name glance

```
openstack service create --name glance --description "OpenStack Image service" image
```

- Create glance service endpoints

```
openstack endpoint create \  
--publicurl http://glance:9292 \  
--internalurl http://glance:9292 \  
--adminurl http://glance:9292 \  
--region RegionOne \  
image
```

- Download an cirros OS image and upload into glance image database
- Login into keystone container and create temporary folder

```
mkdir /tmp/images
```

- Get the cirros image and store it in the temporary folder

```
wget -P /tmp/images http://download.cirros-cloud.net/0.3.4/cirros-0.3.4-x86_64-disk.img
```

Criterion Networks

- Create environment variable file with name openrc_admin1, paste the contents and save it.

```
vi openrc_admin1
```

- Insert the below lines

```
export OS_PROJECT_DOMAIN_ID=default

export OS_USER_DOMAIN_ID=default

export OS_PROJECT_NAME=admin

export OS_TENANT_NAME=admin

export OS_USERNAME=admin

export OS_PASSWORD=admin

export OS_AUTH_URL=http://keystone:35357/v3
```

- Source the openrc_admin1 file

```
source openrc_admin1
```

- Upload image into glance database with name cirros123

```
glance image-create --debug --disk-format qcow2 --container-format bare --file /tmp/images/cirros-0.3.4-x86_64-disk.img --name cirros123
```

- Confirm upload of the images and validate attributes on glance service container

```
glance image-list
```

- Exit the Keystone container

```
exit
```

- Login into the glance container

```
docker ps

docker exec -it <GLANCE-CONTAINERID> /bin/bash
```

NOTE: Also take a note of mysql container ID which will be used within the glance container

- Verify the image is uploaded into glance file directory

Criterion Networks

```
cd /var/lib/glance/images  
ls al
```

(Observe /etc/glance/glance-api.conf)

- Login into Mysql container, login to mysql process and verify available databases

```
mysql -u root -h <mysql-container-id> -p
```

(Provide password as admin123)

```
show databases;
```

- Change to keystone database and list the users

```
use keystone;  
  
select * from user;  
  
exit
```


Kubernetes

Exercise-24

Kubernetes Installation, initialization and join the worker nodes

Kubernetes Installation

NOTE: Perform the following commands on ALL 3 nodes (k8s, knode1 and knode2)

- Update the apt package index:

```
apt-get update
```

- Install the latest version of Docker

```
apt-get install docker.io
```

```
Ubuntu 16.04.4 LTS
Welcome to Ubuntu 16.04.4 LTS (GNU/Linux 4.4.0-116-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage

Get cloud support with Ubuntu Advantage Cloud Guest:
http://www.ubuntu.com/business/services/cloud

54 packages can be updated.
1 update is a security update.

*** System restart required ***
Last login: Fri Jul 20 09:42:47 2018 from 192.168.122.1
root@knode1:~# apt-get update
Get:1 http://security.ubuntu.com/ubuntu xenial-security InRelease [107 kB]
Hit:2 http://archive.ubuntu.com/ubuntu xenial InRelease
Get:3 http://archive.ubuntu.com/ubuntu xenial-updates InRelease [109 kB]
Get:4 http://archive.ubuntu.com/ubuntu xenial-backports InRelease [107 kB]
Fetched 323 kB in 0s (380 kB/s)
Reading package lists... Done
root@knode1:~# apt-get install docker.io
```

- Install packages to allow apt to use a repository over HTTPS:

```
apt-get install -y apt-transport-https curl
```

```
root@knode1:~# apt-get install -y apt-transport-https curl
Reading package lists... Done
Building dependency tree
Reading state information... Done
curl is already the newest version (7.47.0-1ubuntu2.8).
The following packages will be upgraded:
  apt-transport-https
1 upgraded, 0 newly installed, 0 to remove and 52 not upgraded.
Need to get 26.1 kB of archives.
After this operation, 0 B of additional disk space will be used.
Get:1 http://archive.ubuntu.com/ubuntu xenial-updates/main amd64 apt-transport-https amd64 1.2.27 [26.1 kB]
Fetched 26.1 kB in 0s (97.4 kB/s)
(Reading database ... 82138 files and directories currently installed.)
Preparing to unpack .../apt-transport-https_1.2.27_amd64.deb ...
Unpacking apt-transport-https (1.2.27) over (1.2.26) ...
Setting up apt-transport-https (1.2.27) ...
root@knode1:~#
```

- Add Kubernetes official GPG key:

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -
```

```
root@knode1:~# curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add
OK
root@knode1:~#
```

- Add Kubernetes repo to sources.list.d

```
## Multi-Line start

cat <<EOF >/etc/apt/sources.list.d/kubernetes.list

deb http://apt.kubernetes.io/ kubernetes-xenial main

EOF

## Multi-Line end
```

```
root@knode1:~# ## Multi-Line start
root@knode1:~# cat <<EOF >/etc/apt/sources.list.d/kubernetes.list
> deb http://apt.kubernetes.io/ kubernetes-xenial main
> EOF
root@knode1:~# ## Multi-Line end
root@knode1:~#
```

- Update the apt package index and install the kubelet, kubeadm and kubectl packages:

```
apt-get update

apt-get install -y kubelet=1.11.0-00 kubeadm=1.11.0-00 kubectl=1.11.0-00
Correct with: apt-get install -y kubelet kubeadm kubectl
Add: swapoff -a
#nano /etc/fstab
Then comment the /swapfile --> #/swapfile
```

```
root@knode1:~# apt-get update
Hit:1 http://security.ubuntu.com/ubuntu xenial-security InRelease
Hit:3 http://archive.ubuntu.com/ubuntu xenial InRelease
Hit:4 http://archive.ubuntu.com/ubuntu xenial-updates InRelease
Hit:5 http://archive.ubuntu.com/ubuntu xenial-backports InRelease
Get:2 https://packages.cloud.google.com/apt kubernetes-xenial InRelease [8,993 B]
Ign:6 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 Packages
Get:6 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 Packages [18.3 kB]
Fetched 27.3 kB in 0s (31.6 kB/s)
Reading package lists... Done
root@knode1:~# apt-get install -y kubelet=1.11.0-00 kubeadm=1.11.0-00 kubectl=1.11.0-00
```

Kubernetes Master Initialization ONLY on k8s node

- On **k8s node**, initialize the kubernetes master

```
kubeadm init --kubernetes-version=1.11.0 --ignore-preflight-errors=all
```

```
Your Kubernetes master has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

You can now join any number of machines by running the following on each node
as root:

kubeadm join 192.168.122.251:6443 --token 2dj538.0etc9c0wd6v7jtst --discovery-token-ca-cert-hash sha256:562cbc2a625e5665c24ab4de773c09fd0429f493e9683be3734c9ca8be61671

root@k8s1:~#
```

NOTE: Please take a note of the command which is the output of the above command, as that will be used later for joining worker nodes with master. It will be of the format:

```
kubeadm join --token <token> <master-ip>:<master-port> --discovery-token-ca-cert-hash sha256:<hash>
```

- Prepare system for adding workloads, including the network plugin.

```
mkdir -p $HOME/.kube

cp -i /etc/kubernetes/admin.conf $HOME/.kube/config

chown $(id -u):$(id -g) $HOME/.kube/config
```

```
root@k8s1:~# mkdir -p $HOME/.kube
root@k8s1:~# cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
root@k8s1:~# chown $(id -u):$(id -g) $HOME/.kube/config
root@k8s1:~#
```

- Install the Calico network plugin:

```
kubectl apply -f https://docs.projectcalico.org/v2.6/getting-started/kubernet
s/installation/hosted/kubeadm/1.6/calico.yaml
```

pkumar@k8s-master:~\$ kubectl apply -f https://docs.projectcalico.org/v3.14/manifests/calico.yaml

```
root@k8s1:~# kubectl apply -f https://docs.projectcalico.org/v2.6/getting-started/kubernet
es/installation/hosted/kubeadm/1.6/calico.yaml
configmap/calico-config created
daemonset.extensions/calico-etcd created
service/calico-etcd created
daemonset.extensions/calico-node created
deployment.extensions/calico-kube-controllers created
deployment.extensions/calico-policy-controller created
clusterrolebinding.rbac.authorization.k8s.io/calico-cni-plugin created
clusterrole.rbac.authorization.k8s.io/calico-cni-plugin created
serviceaccount/calico-cni-plugin created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
root@k8s1:~#
```

- Check to see if the pods are running:

```
kubectl get pods --all-namespaces
```

```
root@k8s1:~# kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	calico-etcd-lwm2v	1/1	Running	0	39s
kube-system	calico-kube-controllers-74b888b647-vjlgf	0/1	Pending	0	37s
kube-system	calico-node-7cksw	1/2	Running	0	37s
kube-system	coredns-78fcd6894-gcqqb	0/1	Pending	0	2m
kube-system	coredns-78fcd6894-hmskh	0/1	Pending	0	2m
kube-system	etcd-k8s1	1/1	Running	0	1m
kube-system	kube-apiserver-k8s1	1/1	Running	0	1m
kube-system	kube-controller-manager-k8s1	1/1	Running	0	1m
kube-system	kube-proxy-6wxzg	1/1	Running	0	2m
kube-system	kube-scheduler-k8s1	1/1	Running	0	1m

```
root@k8s1:~#
```

(The pods will start up over a short period of time)

Install Kubernetes dashboard plugin

Before joining the nodes, install the kubernetes dashboard, so that it comes on master node.

Deploying the dashboard

```
kubectl create -f https://raw.githubusercontent.com/kubernetes/dashboard/master/aio/d
eploy/recommended/kubernetes-dashboard.yaml
```

If this link doesn't work, try with:

\$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/v2.0.4/aio/deploy/recommended.yaml

Source: <https://github.com/kubernetes/dashboard>

```
root@k8s1:~# kubectl create -f https://raw.githubusercontent.com/kubernetes/dashboard/master/aio/deploy/recommended/kubernetes-dashboard.yaml
secret/kubernetes-dashboard-certs created
serviceaccount/kubernetes-dashboard created
role.rbac.authorization.k8s.io/kubernetes-dashboard-minimal created
rolebinding.rbac.authorization.k8s.io/kubernetes-dashboard-minimal created
deployment.apps/kubernetes-dashboard created
service/kubernetes-dashboard created
```

Verify that the dashboard pod is in running state and is running on k8s1 node, i.e, master node

```
kubectl get pods -o wide --all-namespaces
```

```
root@k8s1:~# kubectl get pods -o wide --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
kube-system	calico-etcd-p88v8	1/1	Running	0	46s	192.168.122.251	k8s1
kube-system	calico-kube-controllers-f5c76676-h8b6l	1/1	Running	0	45s	192.168.122.251	k8s1
kube-system	calico-node-jr952	2/2	Running	0	44s	192.168.122.251	k8s1
kube-system	coredns-78fcd6894-dbzhc	1/1	Running	0	5m	192.168.166.196	k8s1
kube-system	coredns-78fcd6894-hcxfn	1/1	Running	0	5m	192.168.166.195	k8s1
kube-system	etcd-k8s1	1/1	Running	0	4m	192.168.122.251	k8s1
kube-system	kube-apiserver-k8s1	1/1	Running	0	4m	192.168.122.251	k8s1
kube-system	kube-controller-manager-k8s1	1/1	Running	0	4m	192.168.122.251	k8s1
kube-system	kube-proxy-97hrz	1/1	Running	0	5m	192.168.122.251	k8s1
kube-system	kube-scheduler-k8s1	1/1	Running	0	4m	192.168.122.251	k8s1
kube-system	kubernetes-dashboard-5dd89b9875-qm9qn	1/1	Running	0	8s	192.168.166.197	k8s1

To join the worker nodes into Kubernetes master

In all the Worker nodes, run:

NOTE: On knode1 and knode2, perform following commands to join the Kubernetes cluster

```
swapoff -a
#nano /etc/fstab
Then comment the /swapfile --> #/swapfile
```

- Run the below command (which was the output of kubeadm init during master initialization)

```
kubeadm join --token <token> <master-ip>:<master-port> --discovery-token-ca-cert-hash sha256:<hash>
```

check with: kubectl get nodes

A few seconds later, you should notice this node in the output from **kubectl get nodes** when run on the master (k8s1 node)

```
root@knode1:~# kubeadm join 192.168.122.251:6443 --token 2dj538.0etc9c0wd6v7jtst --discovery-token-ca-cert-hash sha256:562cbc2a625e5665c24ab4de773c09fd0429f493e9d683be3734c9ca8be61671
[preflight] running pre-flight checks
[WARNING RequiredIPVSKernelModulesAvailable]: the IPVS proxier will not be used, because the following required kernel modules are not loaded
: [ip_vs ip_vs_rr ip_vs_wrr ip_vs_sh] or no builtin kernel ipvs support: map[ip_vs:{} ip_vs_rr:{} ip_vs_wrr:{} ip_vs_sh:{} nf_conntrack_ipv4:{}]
you can solve this problem with following methods:
1. Run 'modprobe -- ' to load missing kernel modules;
2. Provide the missing builtin kernel ipvs support

I0722 13:02:46.237449 9944 kernel_validator.go:81] Validating kernel version
I0722 13:02:46.237601 9944 kernel_validator.go:96] Validating kernel config
[discovery] Trying to connect to API Server "192.168.122.251:6443"
[discovery] Created cluster-info discovery client, requesting info from "https://192.168.122.251:6443"
[discovery] Requesting info from "https://192.168.122.251:6443" again to validate TLS against the pinned public key
[discovery] Cluster info signature and contents are valid and TLS certificate validates against pinned roots, will use API Server "192.168.122.251:6443"
[discovery] Successfully established connection with API Server "192.168.122.251:6443"
[kubelet] Downloading configuration for the kubelet from the "kubelet-config-1.11" ConfigMap in the kube-system namespace
[kubelet] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[preflight] Activating the kubelet service
[tlsoptions] Waiting for the kubelet to perform the TLS Bootstrap...
[patchnode] Uploading the CRD Socket information "/var/run/dockerhim.sock" to the Node API object "knode1" as an annotation

This node has joined the cluster:
* Certificate signing request was sent to master and a response
  was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the master to see this node join the cluster.
root@knode1:~#
```

Criterion Networks

- Extra commands for help (Do NOT delete node): To delete the node:

```
kubect1 drain <node name> --delete-local-data --force --ignore-daemonsets
```

```
kubect1 delete node <node name>
```

To access the Dashboard:

* <http://localhost:8001/api/v1/namespaces/kubernetes-dashboard/services/https:kubernetes-dashboard:/proxy/> (verify the port 8001 with the kubect1 proxy command)

* `kubect1 create serviceaccount dashboard -n default`

* run:

`kubect1 clusterbinding dashboard-admin -n default \`

`--clusterrole=cluster-admin \`

`--serviceaccount=default:dashboard`

Exercise-25

Working with Pods

- Create a directory in k8s node and a file pod.yml

```
mkdir -p learning  
cd learning
```

```
root@k8s1:~# mkdir -p learning  
root@k8s1:~# cd learning  
root@k8s1:~/learning#
```

- Create a file named pod.yml

```
vim pod.yml
```

- Insert the below lines"

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: myapp-pod  
  labels:  
    app: myapp  
spec:  
  containers:  
    - name: myapp-container  
      image: busybox  
      command: ['sh', '-c', 'echo Hello Kubernetes! && sleep 3600']
```

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  containers:
  - name: myapp-container
    image: busybox
    command: ['sh', '-c', 'echo Hello Kubernetes! && sleep 3600']
```

- Create the pod with specified file

```
kubectl apply -f pod.yml
```

```
root@k8s1:~/learning# kubectl apply -f pod.yml
pod/myapp-pod created
root@k8s1:~/learning#
```

- To introspect pod information

```
kubectl get pods -o wide
```

```
root@k8s1:~/learning# kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE      IP             NODE
myapp-pod     1/1     Running   0           24s      192.168.69.193 knode2
root@k8s1:~/learning#
```

- Another useful command to check pod details

```
kubectl describe pods myapp-pod
```



```

Host Port:      <none>
Command:
  sh
  -c
  echo Hello Kubernetes! && sleep 3600
State:          Running
  Started:      Sun, 22 Jul 2018 13:07:32 +0000
Ready:          True
Restart Count:  0
Environment:    <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-6sltc (ro)
Conditions:
  Type           Status
  Initialized     True
  Ready           True
  ContainersReady True
  PodScheduled    True
Volumes:
  default-token-6sltc:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-6sltc
    Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type     Reason      Age   From          Message
  ----     -
  Normal   Scheduled   49s   default-scheduler   Successfully assigned default/myapp-pod to knode2
  Normal   Pulling     48s   kubelet, knode2     pulling image "busybox"
  Normal   Pulled      47s   kubelet, knode2     Successfully pulled image "busybox"
  Normal   Created     47s   kubelet, knode2     Created container
  Normal   Started     47s   kubelet, knode2     Started container
root@k8s1:~/learning#

```

- To enter into container

```
kubectl exec -it myapp-pod sh
```

```

root@k8s1:~/learning# kubectl exec -it myapp-pod sh
/ #

```

Exercise-26

Working with ReplicaSets

- Create a directory in k8s node and a file replica.yml

- `mkdir /root/learning`
- `cd /root/learning`

```
vim replica.yml
```

```
root@k8s1:~# mkdir /root/learning
root@k8s1:~# cd /root/learning
root@k8s1:/root/learning# vim replica.yml
root@k8s1:/root/learning#
```

- Insert the below lines

```
apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2

kind: ReplicaSet

metadata:

  name: frontend

  labels:

    app: guestbook

    tier: frontend

spec:

  # this replicas value is default

  # modify it according to your case

  replicas: 4

  selector:

    matchLabels:

      tier: frontend
```

```
matchExpressions:
  - {key: tier, operator: In, values: [frontend]}
template:
  metadata:
    labels:
      app: guestbook
      tier: frontend
  spec:
    containers:
      - name: php-redis
        image: gcr.io/google_samples/gb-frontend:v3
        resources:
          requests:
            cpu: 100m
            memory: 100Mi
        env:
          - name: GET_HOSTS_FROM
            value: dns
            # If your cluster config does not include a dns service, then to
            # instead access environment variables to find service host
            # info, comment out the 'value: dns' line above, and uncomment the
            # line below.
            # value: env
    ports:
```

```
- containerPort: 80
```

```
labels:
  app: guestbook
  tier: frontend
spec:
  # this replicas value is default
  # modify it according to your case
  replicas: 4
  selector:
    matchLabels:
      tier: frontend
    matchExpressions:
      - {key: tier, operator: In, values: [frontend]}
  template:
    metadata:
      labels:
        app: guestbook
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
          resources:
            requests:
              cpu: 100m
              memory: 100Mi
          env:
            - name: GET_HOSTS_FROM
              value: dns
              # If your cluster config does not include a dns service, then to
              # instead access environment variables to find service host
              # info, comment out the 'value: dns' line above, and uncomment the
              # line below.
              # value: env
          ports:
            - containerPort: 80
```

- Create the pod with specified file

```
kubectl apply -f replica.yml
```

```
root@k8s1:/root/learning# kubectl apply -f replica.yml
replicaset.apps/frontend created
root@k8s1:/root/learning#
```

- To introspect pod information

```
kubectl get rs -o json
```

Criterion Networks

```
root@k8s1:/root/learning# kubectl describe rs frontend
Name: frontend
Namespace: default
Selector: tier=frontend,tier in (frontend)
Labels: app=guestbook
        tier=frontend
Annotations: kubectl.kubernetes.io/last-applied-configuration={"apiVersion":"apps/v1","kind":"ReplicaSet","metadata":{"annotations":{"ap
p":"guestbook","tier":"frontend"},"name":"frontend","namespace":"...
Replicas: 4 current / 4 desired
Pods Status: 0 Running / 4 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels: app=guestbook
        tier=frontend
  Containers:
    php-redis:
      Image: gcr.io/google_samples/gb-frontend:v3
      Port: 80/TCP
      Host Port: 0/TCP
      Requests:
        cpu: 100m
        memory: 100Mi
      Environment:
        GET_HOSTS_FROM: dns
      Mounts:
        <none>
      Volumes:
        <none>
Events:
  Type      Reason              Age             From              Message
  ----      -
  Normal    SuccessfulCreate    33s             replicaset-controller Created pod: frontend-qzrpv
  Normal    SuccessfulCreate    33s             replicaset-controller Created pod: frontend-rn8dh
  Normal    SuccessfulCreate    33s             replicaset-controller Created pod: frontend-zn2lh
  Normal    SuccessfulCreate    33s             replicaset-controller Created pod: frontend-2dfbk
root@k8s1:/root/learning#
```

- Another useful command to check pod details

```
kubectl describe rs frontend
```

```
root@k8s1:/root/learning# kubectl get pods -o wide
NAME          READY   STATUS             RESTARTS   AGE   IP           NODE
frontend-2dfbk 0/1     ContainerCreating   0          1m    <none>       knode2
frontend-qzrpv 0/1     ContainerCreating   0          1m    <none>       knode1
frontend-rn8dh 0/1     ContainerCreating   0          1m    <none>       knode1
frontend-zn2lh 0/1     ContainerCreating   0          1m    <none>       knode2
myapp-pod      1/1     Running             0          4m    192.168.69.193 knode2
root@k8s1:/root/learning#
```

- Display all the pod information including their IP addresses

```
kubectl get pods -o wide
```

```
root@k8s1:/root/learning# kubectl get pods -l tier=frontend
NAME          READY   STATUS    RESTARTS   AGE
frontend-2dfbk 1/1     Running   0          1m
frontend-qzrpv 1/1     Running   0          1m
frontend-rn8dh 1/1     Running   0          1m
frontend-zn2lh 1/1     Running   0          1m
root@k8s1:/root/learning#
```

- Display selected pod information with labels

```
kubectl get pods -l tier=frontend
```

```
root@k8s1:/root/learning# cd /root/learning
root@k8s1:/root/learning# vim autoscale.yml
```

- Autoscale options with ReplicaSets

```
cd /root/learning
```

```
vim autoscale.yml
```

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: frontend-scaler
spec:
  scaleTargetRef:
    kind: ReplicaSet
    name: frontend
  minReplicas: 3
  maxReplicas: 10
  targetCPUUtilizationPercentage: 50
~
~
~
```

- Insert the below lines

```
apiVersion: autoscaling/v1

kind: HorizontalPodAutoscaler

metadata:

  name: frontend-scaler

spec:

  scaleTargetRef:

    kind: ReplicaSet

    name: frontend

  minReplicas: 3

  maxReplicas: 10

  targetCPUUtilizationPercentage: 50
```

```
root@k8s1:/root/learning# kubectl apply -f autoscale.yml
horizontalpodautoscaler.autoscaling/frontend-scaler created
root@k8s1:/root/learning#
```

- Apply the autoscale file

```
kubectl apply -f autoscale.yml
```

```
root@k8s1:/root/learning# kubectl delete rs/frontend --cascade=false
replicaset.extensions "frontend" deleted
root@k8s1:/root/learning# kubectl get rs
No resources found.
root@k8s1:/root/learning# kubectl get pods

```

NAME	READY	STATUS	RESTARTS	AGE
frontend-2dfbk	1/1	Running	0	3m
frontend-qzrpv	1/1	Running	0	3m
frontend-rn8dh	1/1	Running	0	3m
frontend-zn2lh	1/1	Running	0	3m
myapp-pod	1/1	Running	0	6m

```
root@k8s1:/root/learning#
```

- Cascading option (Demonstrates loosely couple nature of replicaset and pods)

```
kubectl delete rs/frontend --cascade=false
```

```
kubectl get rs
```

```
kubectl get pods
```

Exercise-27

Working with Deployments

- Create a new file in k8s node with name deploy1.yml

```
cd /root/learning/  
vim deploy1.yml
```

- Insert the below lines

```
apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2  
kind: Deployment  
metadata:  
  name: nginx-deployment  
  labels:  
    app: nginx  
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      app: nginx  
  template:  
    metadata:  
      labels:  
        app: nginx  
    spec:  
      containers:  
        - name: nginx
```



```
image: nginx:1.7.9

ports:

- containerPort: 80
```

```
apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
        ports:
        - containerPort: 80
~
~
```

- Create the pod with specified file

```
kubectl apply -f deploy1.yml --record
```

```
root@k8s1:/root/learning# kubectl apply -f deploy1.yml --record
deployment.apps/nginx-deployment created
root@k8s1:/root/learning#
```

- To introspect pod information

```
kubectl get deployments
```

```
root@k8s1:/root/learning# kubectl get deployments
NAME                DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    3         3         3            1          23s
root@k8s1:/root/learning#
```

- Another useful command to check pod details

```
kubectl describe deployments nginx-deployment
```

```
root@k8s1:/root/learning# kubectl describe deployments nginx-deployment
Name:          nginx-deployment
Namespace:     default
CreationTimestamp: Sun, 22 Jul 2018 13:16:22 +0000
Labels:        app=nginx
Annotations:    deployment.kubernetes.io/revision=1
                kubectl.kubernetes.io/last-applied-configuration={"apiVersion":"apps/v1","kind":"Deployment","metadata":{"annotations":{"kubernetes.io/change-cause":"kubectl apply --filename=deploy1.yml --record=true...
                kubernetes.io/change-cause=kubectl apply --filename=deploy1.yml --record=true
Selector:      app=nginx
Replicas:      3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType:  RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=nginx
  Containers:
    nginx:
      Image:      nginx:1.7.9
      Port:       80/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:      <none>
      Volumes:     <none>
  Conditions:
    Type           Status  Reason
    ----           -
    Available      True    MinimumReplicasAvailable
    Progressing    True    NewReplicaSetAvailable
    OldReplicaSets: <none>
    NewReplicaSet:  nginx-deployment-67594d6bf6 (3/3 replicas created)
  Events:
    Type     Reason              Age   From              Message
    ----     -
    Normal   ScalingReplicaSet   47s   deployment-controller   Scaled up replica set nginx-deployment-67594d6bf6 to 3
root@k8s1:/root/learning#
```

- Display all the pod information along with labels

```
kubectl get pods --show-labels
```

```
root@k8s1:/root/learning# kubectl get pods --show-labels
NAME                                READY    STATUS    RESTARTS   AGE    LABELS
frontend-2dfbk                     1/1     Running   0           6m     app=guestbook,tier=frontend
frontend-qzrpv                     1/1     Running   0           6m     app=guestbook,tier=frontend
frontend-rn8dh                     1/1     Running   0           6m     app=guestbook,tier=frontend
frontend-zn2lh                     1/1     Running   0           6m     app=guestbook,tier=frontend
myapp-pod                          1/1     Running   0          10m     app=myapp
nginx-deployment-67594d6bf6-8n5fw  1/1     Running   0           1m     app=nginx,pod-template-hash=2315082692
nginx-deployment-67594d6bf6-cn46g  1/1     Running   0           1m     app=nginx,pod-template-hash=2315082692
nginx-deployment-67594d6bf6-qdgph  1/1     Running   0           1m     app=nginx,pod-template-hash=2315082692
root@k8s1:/root/learning#
```

- Display selected pod information with labels

```
kubectl get pods -l app=nginx
```

```
root@k8s1:/root/learning# kubectl get pods -l app=nginx
NAME                                READY    STATUS    RESTARTS   AGE
nginx-deployment-67594d6bf6-8n5fw  1/1     Running   0           1m
nginx-deployment-67594d6bf6-cn46g  1/1     Running   0           1m
nginx-deployment-67594d6bf6-qdgph  1/1     Running   0           1m
root@k8s1:/root/learning#
```

- To check the roll-out status

```
kubectl rollout status deployment/nginx-deployment
```

```
root@k8s1:/root/learning# kubectl rollout status deployment/nginx-deployment
deployment "nginx-deployment" successfully rolled out
root@k8s1:/root/learning#
```

- Check the ReplicaSet controllers

```
kubectl get rs
```

```
root@k8s1:/root/learning# kubectl get rs
NAME                                DESIRED    CURRENT    READY    AGE
nginx-deployment-67594d6bf6         3          3          3        2m
root@k8s1:/root/learning#
```

- To rollout an update to nginx version

```
kubectl set image deployment/nginx-deployment nginx=nginx:1.9.1
```

```
root@k8s1:/root/learning# kubectl set image deployment/nginx-deployment nginx=nginx:1.9.1
deployment.extensions/nginx-deployment image updated
root@k8s1:/root/learning#
```

- Check the ReplicaSet controllers again

```
kubectl get rs
```

```
root@k8s1:/root/learning# kubectl get rs
NAME                                DESIRED    CURRENT    READY    AGE
nginx-deployment-67594d6bf6         2          2          2        3m
nginx-deployment-6fdbb596db         2          2          1        35s
root@k8s1:/root/learning#
```

- Check the pods and verify it shows new pods

```
kubectl get pods
```

```
root@k8s1:/root/learning# kubectl get pods
NAME                                READY      STATUS    RESTARTS   AGE
frontend-2dfbk                     1/1        Running   0           9m
frontend-qzrpv                      1/1        Running   0           9m
frontend-rn8dh                      1/1        Running   0           9m
frontend-zn2lh                      1/1        Running   0           9m
myapp-pod                          1/1        Running   0          13m
nginx-deployment-6fdbb596db-mxqdm   1/1        Running   0           41s
nginx-deployment-6fdbb596db-w78dl   1/1        Running   0           1m
nginx-deployment-6fdbb596db-wsxcw   1/1        Running   0          22s
root@k8s1:/root/learning#
```

Criterion Networks

- Check the deployments workflow

```
kubectl describe deployments
```

- Check the rollout history of the deployments

```
kubectl rollout history deployment/nginx-deployment
```

```
root@k8s1:/root/learning# kubectl rollout history deployment/nginx-deployment
deployments "nginx-deployment"
REVISION  CHANGE-CAUSE
1         kubectl apply --filename=deploy1.yml --record=true
2         kubectl apply --filename=deploy1.yml --record=true

root@k8s1:/root/learning#
```

- To change the revision limit

```
kubectl patch deployment/nginx-deployment -p '{"spec":{"revisionHistoryLimit":10}}'
```

```
root@k8s1:/root/learning# kubectl patch deployment/nginx-deployment -p '{"spec":{"revisionHistoryLimit":10}}'
deployment.extensions/nginx-deployment not patched
root@k8s1:/root/learning# kubectl autoscale deployment nginx-deployment --min=10 --max=15 --cpu-percent=80
horizontalpodautoscaler.autoscaling/nginx-deployment autoscaled
root@k8s1:/root/learning#
```

- Scaling a deployment

```
kubectl autoscale deployment nginx-deployment --min=10 --max=15 --cpu-percent=80
```

```
root@k8s1:/root/learning# kubectl rollout pause deployment/nginx-deployment
deployment.extensions/nginx-deployment paused
root@k8s1:/root/learning#
```

- Pause a deployment

```
kubectl rollout pause deployment/nginx-deployment
```

- Change the image version names

```
kubectl set image deploy/nginx-deployment nginx=nginx:1.9.4
```

- Resume a deployment

```
kubectl rollout resume deployment/nginx-deployment
```

- Waiting time before declaring as dead

```
kubectl patch deployment/nginx-deployment -p '{"spec":{"progressDeadlineSeconds":600}}'
```

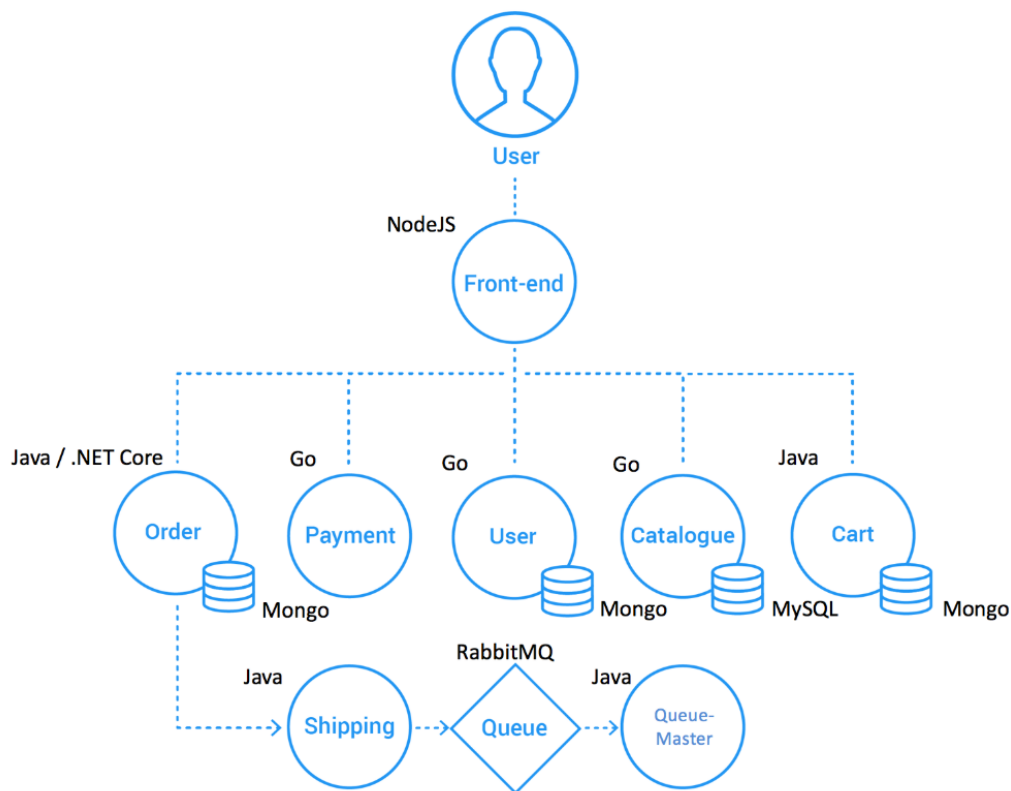
CONFIDENTIAL

Exercise-28

Objective

In this set of exercises, we will deploy a simple multi-tier application name sock-shop on the kubernetes cluster

Sock-Shop Architecture Overview



Sock Shop is a multi-tier e-commerce site which is designed to test container based environments. All the components of sock-shop are designed to be containers which can be deployed in our kubernetes cluster.

Clone Sock Shop Deployment Repo

On **k8s1** node

Clone the repository:

```
cd /home/ubuntu/
git clone https://github.com/microservices-demo/microservices-demo.git
```

```
ubuntu@k8s:~$ sudo su
sudo: unable to resolve host k8s
root@k8s:/home/ubuntu# git clone https://github.com/microservices-demo/microservices-demo
Cloning into 'microservices-demo'...
remote: Counting objects: 9603, done.
remote: Compressing objects: 100% (6/6), done.
remote: Total 9603 (delta 1), reused 1 (delta 1), pack-reused 9596
Receiving objects: 100% (9603/9603), 52.83 MiB | 39.80 MiB/s, done.
Resolving deltas: 100% (5770/5770), done.
Checking connectivity... done.
root@k8s:/home/ubuntu#
```

Go to the repository:

```
cd microservices-demo

ls
```

```
root@k8s:/home/ubuntu# cd microservices-demo/
root@k8s:/home/ubuntu/microservices-demo# ls
deploy  graphs      install      LICENSE      push.sh      shippable.jobs.yml      shippable.triggers.yml
docs    healthcheck internal-docs openapi      README.md    shippable.resources.yml staging
```

Create namespace and deploy application

- Clone the repository

```
cd /home/ubuntu/

git clone https://github.com/microservices-demo/microservices-demo.git
```

- Go to the repository:

```
cd /home/ubuntu/microservices-demo

ls
```

- Create new namespace sock-shop:

```
kubectl create namespace sock-shop
```

- To verify and list all the namespaces:

```
kubectl get namespaces
```

```
root@k8s1:~# cd microservices-demo/
root@k8s1:~/microservices-demo# ls
deploy  graphs      install      LICENSE  push.sh    shippable.jobs.yml    shippable.triggers.yml
docs    healthcheck internal-docs openapi   README.md  shippable.resources.yml staging
root@k8s1:~/microservices-demo# kubectl create namespace sock-shop
namespace/sock-shop created
root@k8s1:~/microservices-demo# kubectl get namespaces
NAME              STATUS    AGE
default           Active   16m
kube-public       Active   16m
kube-system       Active   16m
sock-shop         Active   13s
root@k8s1:~/microservices-demo#
```

- Deploy sock-shop application:

```
cd /home/ubuntu/microservices-demo/deploy/kubernetes/
kubectl apply -f complete-demo.yaml -n sock-shop
```

```
root@k8s1:~/microservices-demo/deploy/kubernetes# kubectl apply -f complete-demo.yaml -n sock-shop
deployment.extensions/carts-db created
service/carts-db created
deployment.extensions/carts created
service/carts created
deployment.extensions/catalogue-db created
service/catalogue-db created
deployment.extensions/catalogue created
service/catalogue created
deployment.extensions/front-end created
service/front-end created
deployment.extensions/orders-db created
service/orders-db created
deployment.extensions/orders created
service/orders created
deployment.extensions/payment created
service/payment created
deployment.extensions/queue-master created
service/queue-master created
deployment.extensions/rabbitmq created
service/rabbitmq created
deployment.extensions/shipping created
service/shipping created
deployment.extensions/user-db created
service/user-db created
deployment.extensions/user created
service/user created
root@k8s1:~/microservices-demo/deploy/kubernetes#
```

Verify Pods Status

- To get the details of pods running in sock-shop namespace:

```
kubectl get pods -n sock-shop
```



```

root@k8s1:~/microservices-demo/deploy/kubernetes# kubectl get pods -n sock-shop
NAME                                READY    STATUS              RESTARTS   AGE
carts-6dfdc59f8-sx6rw              1/1     Running             0          38s
carts-db-6c9b649b49-ts9v           0/1     ContainerCreating   0          39s
catalogue-7d7f9f87f-59szv          0/1     ContainerCreating   0          38s
catalogue-db-745c877d4f-kvljb       0/1     ContainerCreating   0          38s
front-end-6f779bdb68-nmjxb         0/1     ContainerCreating   0          38s
orders-5c4f477565-xzcg6            0/1     ContainerCreating   0          36s
orders-db-db498cfb9-srkch          0/1     ContainerCreating   0          37s
payment-5df6dc6bcc-j55bg           0/1     ContainerCreating   0          35s
queue-master-787b68b7fd-sd84x      0/1     ContainerCreating   0          34s
rabbitmq-86fcc47fc-qtnv2           0/1     ContainerCreating   0          33s
shipping-64f8c7558c-72fk4         0/1     ContainerCreating   0          32s
user-7848fb86db-pvrv7              0/1     ContainerCreating   0          31s
user-db-586b8566b4-kdsnh           0/1     ContainerCreating   0          32s
root@k8s1:~/microservices-demo/deploy/kubernetes#

```

[Here we can see that pods are still building.]

Note: Please check the status after 2 to 3 minutes:]

- You can also get the services running over sock-shop namespaces.

```
kubectl get services -n sock-shop -o wide
```

```

root@k8s1:~/microservices-demo/deploy/kubernetes# kubectl get services -n sock-shop -o wide
NAME            TYPE        CLUSTER-IP      EXTERNAL-IP  PORT(S)          AGE    SELECTOR
carts           ClusterIP   10.105.254.99   <none>       80/TCP           28m    name=carts
carts-db        ClusterIP   10.103.148.98   <none>       27017/TCP        28m    name=carts-db
catalogue       ClusterIP   10.107.44.182   <none>       80/TCP           28m    name=catalogue
catalogue-db    ClusterIP   10.109.59.168   <none>       3306/TCP         28m    name=catalogue-db
front-end       NodePort    10.107.156.141   <none>       80:30001/TCP     28m    name=front-end
orders          ClusterIP   10.107.221.115   <none>       80/TCP           28m    name=orders
orders-db       ClusterIP   10.97.46.116     <none>       27017/TCP        28m    name=orders-db
payment         ClusterIP   10.102.233.205   <none>       80/TCP           28m    name=payment
queue-master    ClusterIP   10.106.122.172   <none>       80/TCP           28m    name=queue-master
rabbitmq        ClusterIP   10.107.66.31     <none>       5672/TCP         28m    name=rabbitmq
shipping        ClusterIP   10.97.232.224    <none>       80/TCP           28m    name=shipping
user            ClusterIP   10.110.188.75    <none>       80/TCP           28m    name=user
user-db         ClusterIP   10.106.134.150   <none>       27017/TCP        28m    name=user-db
root@k8s1:~/microservices-demo/deploy/kubernetes#

```

Note: Get the node port for the front-end service as highlighted in below screenshot

- Verify that all the pods came to running state and fetch the node name

```
kubectl get pods -n sock-shop -o wide
```

Criterion Networks

```
root@k8s1:~/microservices-demo/deploy/kubernetes# kubectl get pods -n sock-shop -o wide
NAME                                READY    STATUS    RESTARTS   AGE    IP              NODE
carts-6dfdc59f8-sx6rw               1/1     Running   0           9m     192.168.195.129 knode1
carts-db-6c9b649b49-ts9v            1/1     Running   0           9m     192.168.69.193  knode2
catalogue-7d7f9f87f-59szv           1/1     Running   0           9m     192.168.195.130 knode1
catalogue-db-745c877d4f-kvljb       1/1     Running   0           9m     192.168.69.195  knode2
front-end-6f779bdb68-nmjxb          1/1     Running   0           9m     192.168.69.194  knode2
orders-5c4f477565-xzcg6             1/1     Running   0           9m     192.168.69.196  knode2
orders-db-db498cfb9-srkch           1/1     Running   0           9m     192.168.195.131 knode1
payment-5df6dc6bcc-j55bg            1/1     Running   0           9m     192.168.195.132 knode1
queue-master-787b68b7fd-sd84x       1/1     Running   0           9m     192.168.69.197  knode2
rabbitmq-86fcc47fc-qtnv2            1/1     Running   0           9m     192.168.195.133 knode1
shipping-64f8c7558c-72fk4           1/1     Running   0           9m     192.168.69.198  knode2
user-7848fb86db-pvvnv7              1/1     Running   0           9m     192.168.69.199  knode2
user-db-586b8566b4-kdsnh            1/1     Running   0           9m     192.168.195.134 knode1
root@k8s1:~/microservices-demo/deploy/kubernetes#
```

Note: Get the node name against the front-end pod

- Get the IP of the above captured node

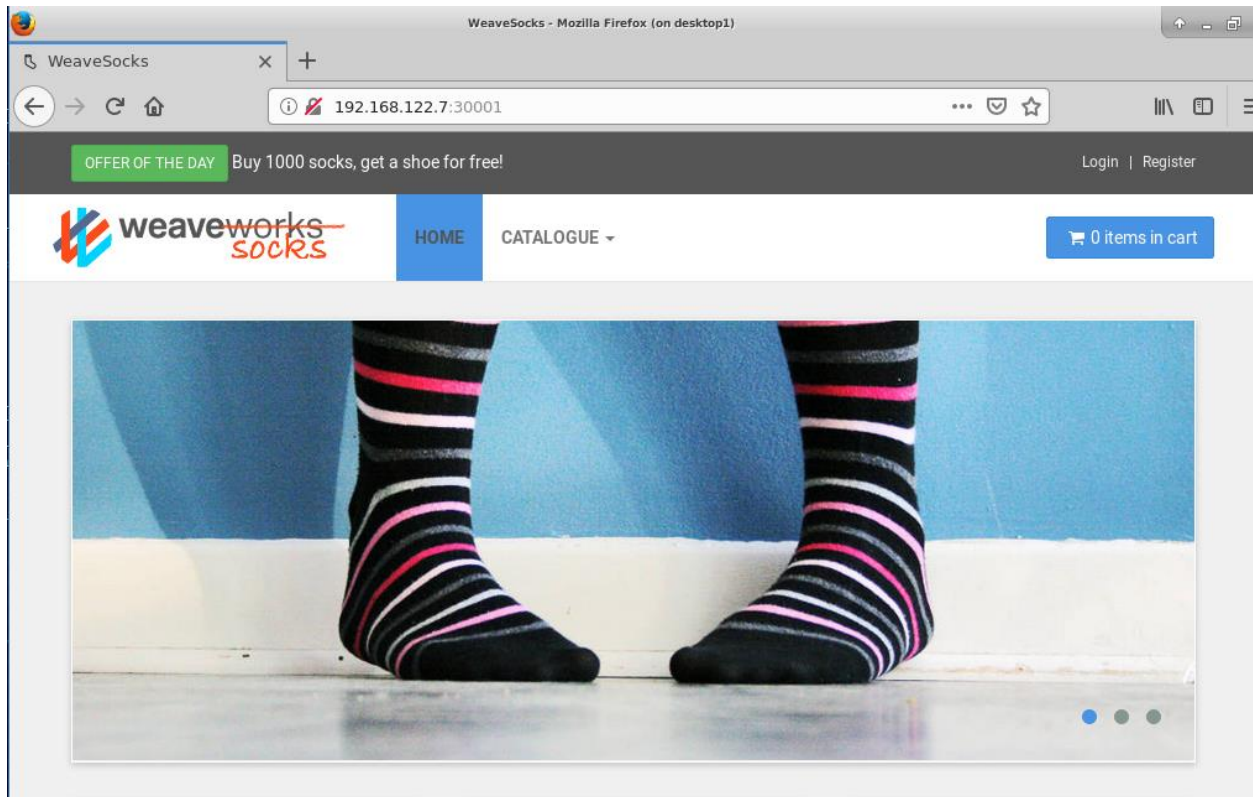
```
kubectl get nodes -o wide
```

```
root@k8s1:~/microservices-demo/deploy/kubernetes# kubectl get nodes -o wide
NAME      STATUS    ROLES    AGE    VERSION    INTERNAL-IP    EXTERNAL-IP    OS-IMAGE             KERNEL-VERSION    CONTAINER-RUNTIME
k8s1     Ready     master   29m    v1.11.0    192.168.122.251 <none>         Ubuntu 16.04.4 LTS  4.4.0-116-generic   docker://18.6.1
knode1   Ready     <none>    19m    v1.11.0    192.168.122.126 <none>         Ubuntu 16.04.4 LTS  4.4.0-116-generic   docker://18.6.1
knode2   Ready     <none>    19m    v1.11.0    192.168.122.7   <none>         Ubuntu 16.04.4 LTS  4.4.0-116-generic   docker://18.6.1
root@k8s1:~/microservices-demo/deploy/kubernetes#
```

- Go to the Desktop VNC session, open the browser and browse the below URL

```
http://<node-ip>:<node-port>
```

Criterion Networks



Exercise-29

Config Map

Problem Statement

Let's us consider the use case where some environment variables such as username and password for a database are to be used in multiple replication controller or pod definition files. The username and password value would need to be specified in each of the definition files and if the username and password were to change, all the definition files would need to be updated as well, which could be very tedious.

Alternatively, variable values could be supplied to kubectl when a command is run, which involves specifying command-line flags each time the command is run.

Solution

Here comes the usage of Kubernetes ConfigMap management pattern which is a map of configuration properties that can be used in definition files for pods, replication controllers, and other Kubernetes objects to configure

- environment variables,
- command arguments, and
- configuration files such as key-value pairs in volumes, etc.

A single ConfigMap may package multiple configuration properties as key/value pairs. By creating ConfigMaps, you specify the configuration properties in a single configuration map, which can be updated as required without having to update each of the definition files in which the ConfigMap is used.

Decoupling the containers from the configuration data provides portability of the applications running in the containers.

Exercise-30

Config map use case

Create config file

The following configuration file will create a configuration file that keeps a list of addresses, open a vim editor with a file name **configmap.yaml**

```
cd /home/ubuntu/learning/  
vim configmap.yaml
```

Then copy and paste the below configuration in the opened vim editor

```
---  
apiVersion: v1  
kind: ConfigMap  
metadata:  
  name: db-config  
  namespace: default  
data:  
  db-ip-addresses: 1.2.3.4,5.6.7.8
```

To create the config use the below command

```
kubectl create -f configmap.yaml
```

Below message will be displayed

configmap/db-config created

The data section contains all the key-value pairs, in this case, just a single pair with a key name of db-ip-addresses. It will be important later when consuming the configmap in a pod.

You can check out the content to make sure configmap is created:

```
kubectl get configmap db-config -o yaml
```

Output will be displayed as below:

```
apiVersion: v1

data:
  db-ip-addresses: 1.2.3.4,5.6.7.8

kind: ConfigMap

metadata:
  creationTimestamp: 2018-11-23T07:32:52Z
  name: db-config
  namespace: default
  resourceVersion: "382141"
  selfLink: /api/v1/namespaces/default/configmaps/db-config
  uid: fc440309-eef1-11e8-b862-5254006dd4e6
```

Consuming a ConfigMap as an environment variable

When we are creating a pod, we can specify a ConfigMap and consume its values.

```
vim mypod1.yaml
```

Go to insert mode (by pressing "i") and copy-paste the below configuration in this file

Here is how to consume our configuration map as an environment variable:

```
apiVersion: v1

kind: Pod

metadata:
  name: some-pod

spec:
  containers:
    - name: some-container
```

```

    image: busybox

    command: [ "/bin/sh", "-c", "env" ]

    env:

      - name: DB_IP_ADDRESSES

        valueFrom:

          configMapKeyRef:

            name: db-config

            key: db-ip-addresses

    restartPolicy: Never

```

Create pod with above config

```
kubectl create -f mypod1.yaml
```

This pod runs the busybox minimal container and executes an env bash command and immediately exists. The db-ip-addresses key from the db-config map is mapped to the DB_IP_ADDRESSES environment variable, and is reflected in the output:

```
kubectl logs some-pod
```

```

root@k8s1:~# kubectl logs some-pod
KUBERNETES_SERVICE_PORT=443
KUBERNETES_PORT=tcp://10.96.0.1:443
HOSTNAME=some-pod
SHLVL=1
HOME=/root
DB_IP_ADDRESSES=1.2.3.4,5.6.7.8
KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
KUBERNETES_PORT_443_TCP_PORT=443
KUBERNETES_PORT_443_TCP_PROTO=tcp
KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443
KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_SERVICE_HOST=10.96.0.1
PWD=/
root@k8s1:~#

```

Exercise-31

Starting the Kubernetes UI application

Starting the Dashboard application

[The below commands has to be executed in k8s1 node]

Get the IP of the k8s1 node using below command

```
ifconfig ens3
```

Mostly the IP will be "192.168.122.251", but verify in your setup once.

```
root@k8s1:~# ifconfig ens3
ens3      Link encap:Ethernet  HWaddr 52:54:00:de:62:6f
          inet addr:192.168.122.251  Bcast:192.168.122.255  Mask:255.255.255.0
          inet6 addr: fe80::5054:ff:fede:626f/64  Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:303946 errors:0 dropped:18 overruns:0 frame:0
          TX packets:116787 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:617553694 (617.5 MB)  TX bytes:22167798 (22.1 MB)

root@k8s1:~#
```

Start the dashboard application using proxy server command

```
kubect1 proxy --address='192.168.122.251' --port=8001 --accept-hosts='^*$' &
```

```
root@k8s1:~# kubect1 proxy --address='192.168.122.251' --port=8001 --accept-hosts='^*$' &
[1] 2826
root@k8s1:~# Starting to serve on 192.168.122.251:8001
```

NOTE: Check if the IP is same as mentioned in above command, else replace the above mentioned IP with the one which you have obtained using the previous command.

As we will be accessing the dashboard from desktop node, so we have to change the service type to **NodePort** from **ClusterIP**

Edit kubernetes-dashboard service

```
kubect1 -n kube-system edit service kubernetes-dashboard
```



```

selfLink: /api/v1/namespaces/kube-system/services/kubernetes-dashboard
uid: 298979d0-17c8-11e9-9e27-525400de626f
spec:
  clusterIP: 10.106.103.18
  ports:
  - port: 443
    protocol: TCP
    targetPort: 8443
  selector:
    k8s-app: kubernetes-dashboard
  sessionAffinity: Nonen.yaml" 27L, 739C
  type: ClusterIP
status:
  loadBalancer: {}

```

Replace the ClusterIP to NodePort

You should see yaml representation of the service. Change type: ClusterIP to type: NodePort and save file. If it's already changed go to next step.

Next we need to check port on which Dashboard was exposed.

```
kubectl -n kube-system get service kubernetes-dashboard
```

```

root@k8s1:~# kubectl -n kube-system get service kubernetes-dashboard
NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes-dashboard NodePort    10.106.103.18 <none>         443:31466/TCP    2h
root@k8s1:~#

```

Dashboard has been exposed on port obtained above (HTTPS). Now you can access it from your browser at: <https://master-ip:nodeport>.

Create service account and cluster role binding

Create cluster admin service account (on *k8s1* node)

```

kubectl create serviceaccount cluster-admin-dashboard-sa

kubectl create clusterrolebinding cluster-admin-dashboard-sa \

    --clusterrole=cluster-admin \

    --serviceaccount=default:cluster-admin-dashboard-sa

```

And then, use the token of just created cluster admin service account.

```

kubectl get secret | grep cluster-admin-dashboard-sa

kubectl describe secret <cluster-admin-dashboard-name-obtained-above>

```

Copy this output and open the **desktop node** from the "Access Devices" tab and create a file named token using vim editor

Go to the insert mode and paste the token here

CONFIDENTIAL

Exercise-32

Kubernetes Dashboard

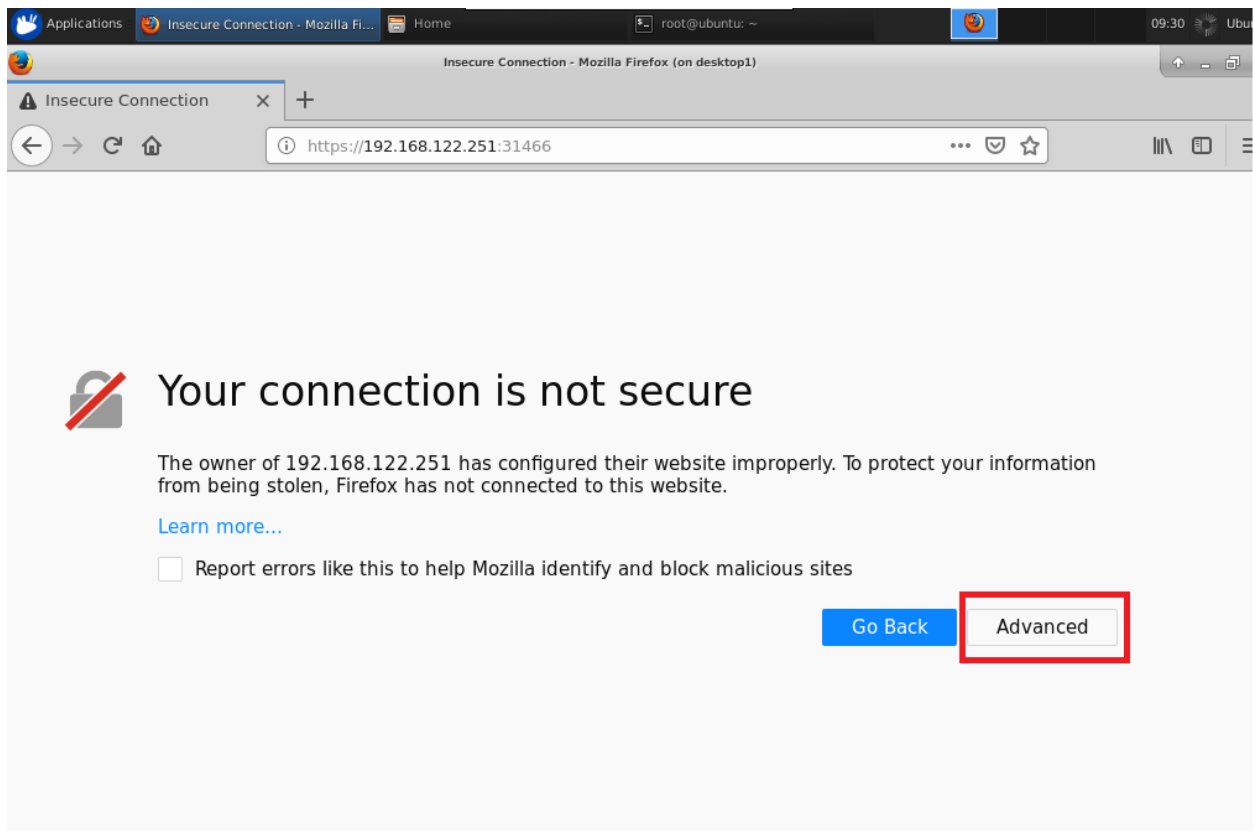
Verify the Dashboard

Go to the desktop's VNC session and open the Mozilla firefox browser and type the below URL

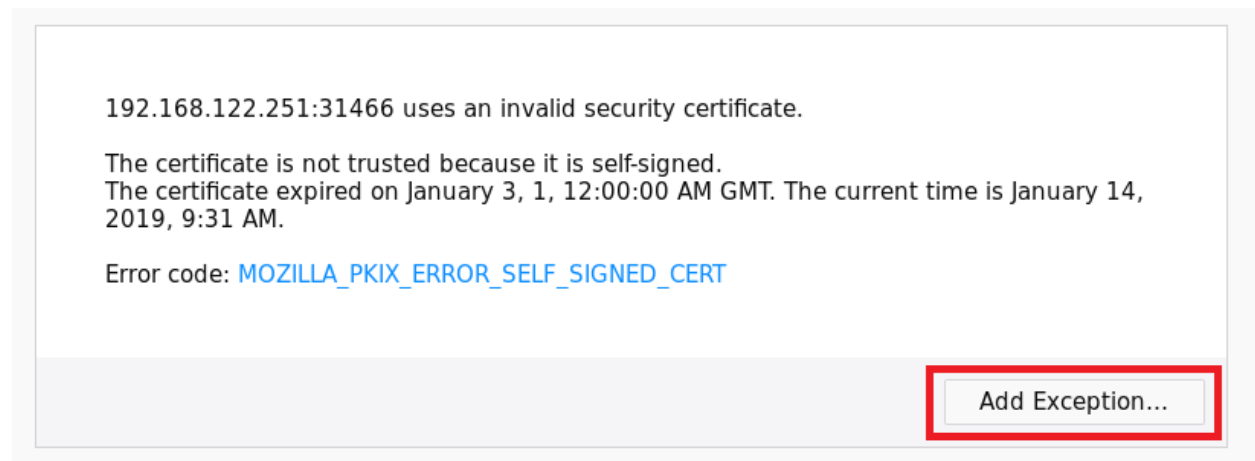
```
https://<k8s1-node-ip>:<kubernetes-dashboard-node-port>
```

NOTE: Kubernetes dashboard node port is the port we have obtained previously.

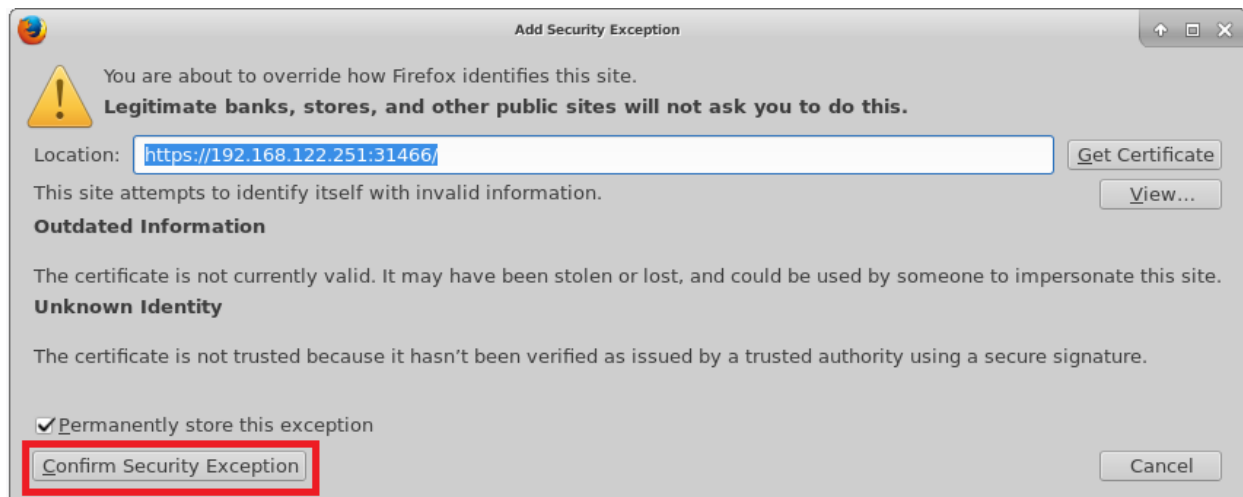
It will redirect to the below page, click on advanced



Scroll down and click on "Add Exception"

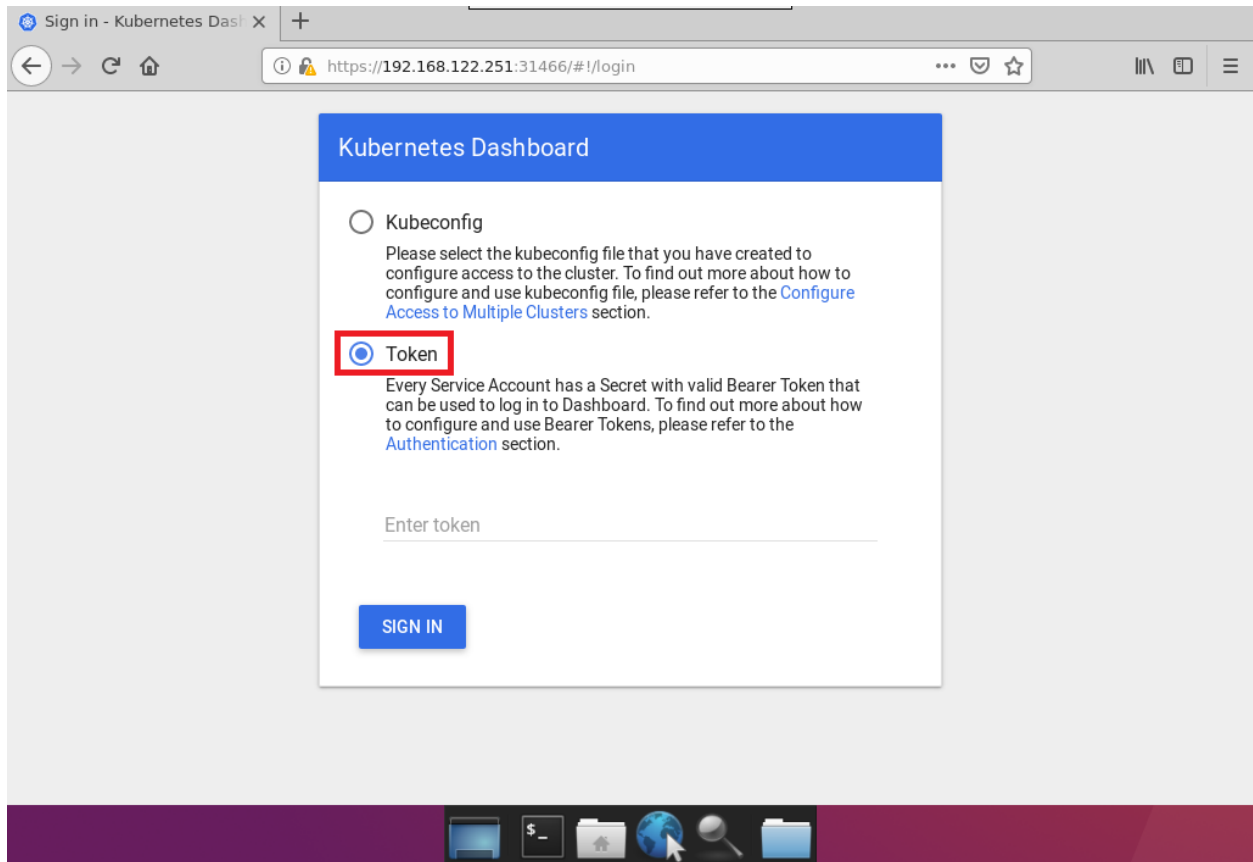


A window as shown below will pop-up click on "Confirm Security Exception"




The Kubernetes Dashboard's authentication page will open, there select "Token"

Criterion Networks

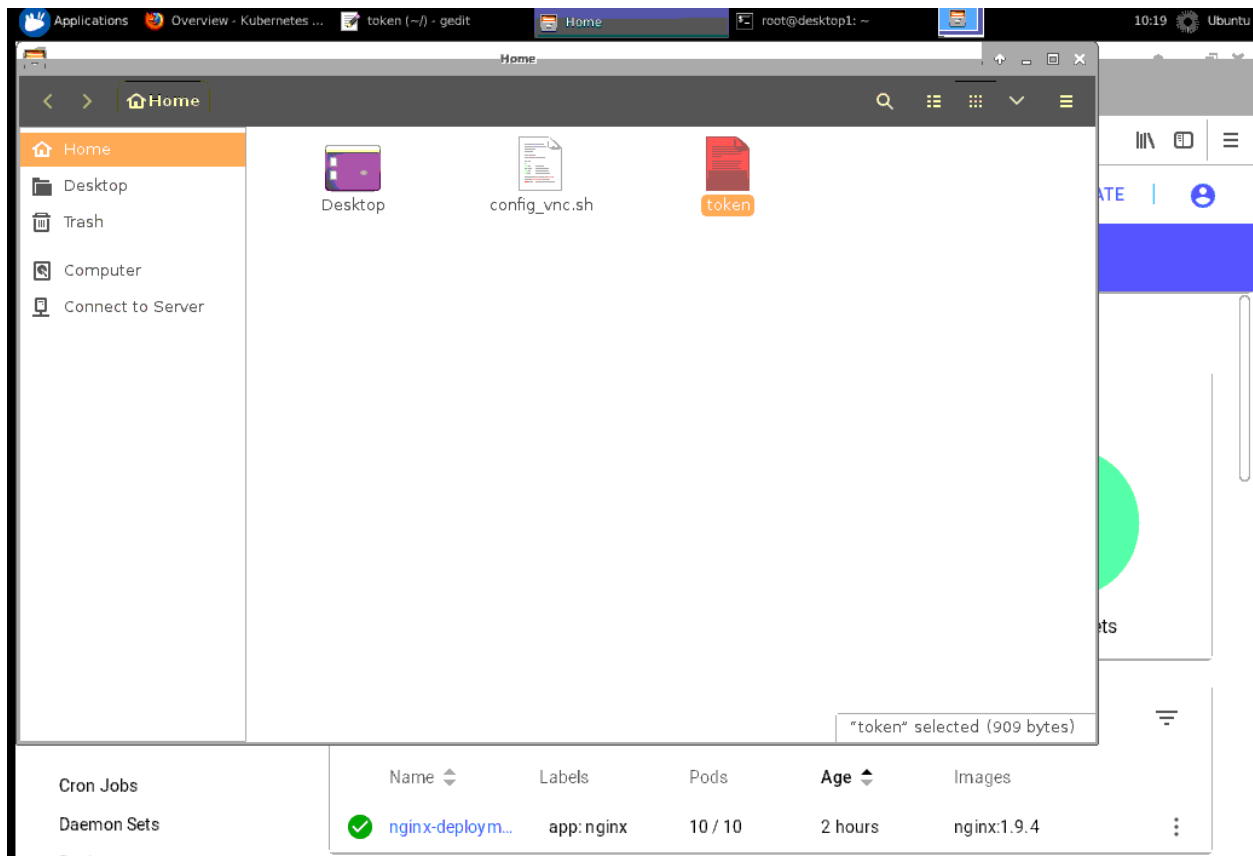


Now in the VNC session open the terminal and read the token file which was created earlier in this exercise. Double click on the token file to open it.

Criterion Networks



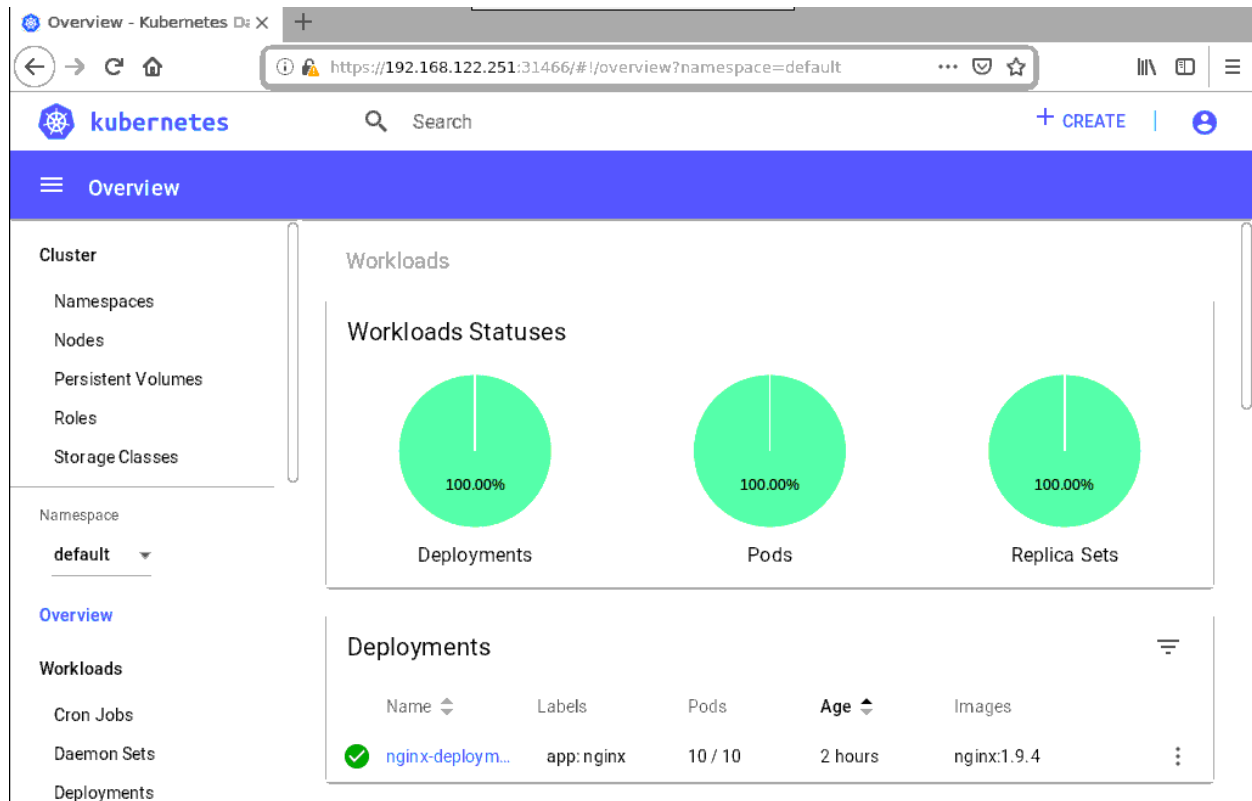
Copy this token and paste in the Dashboard



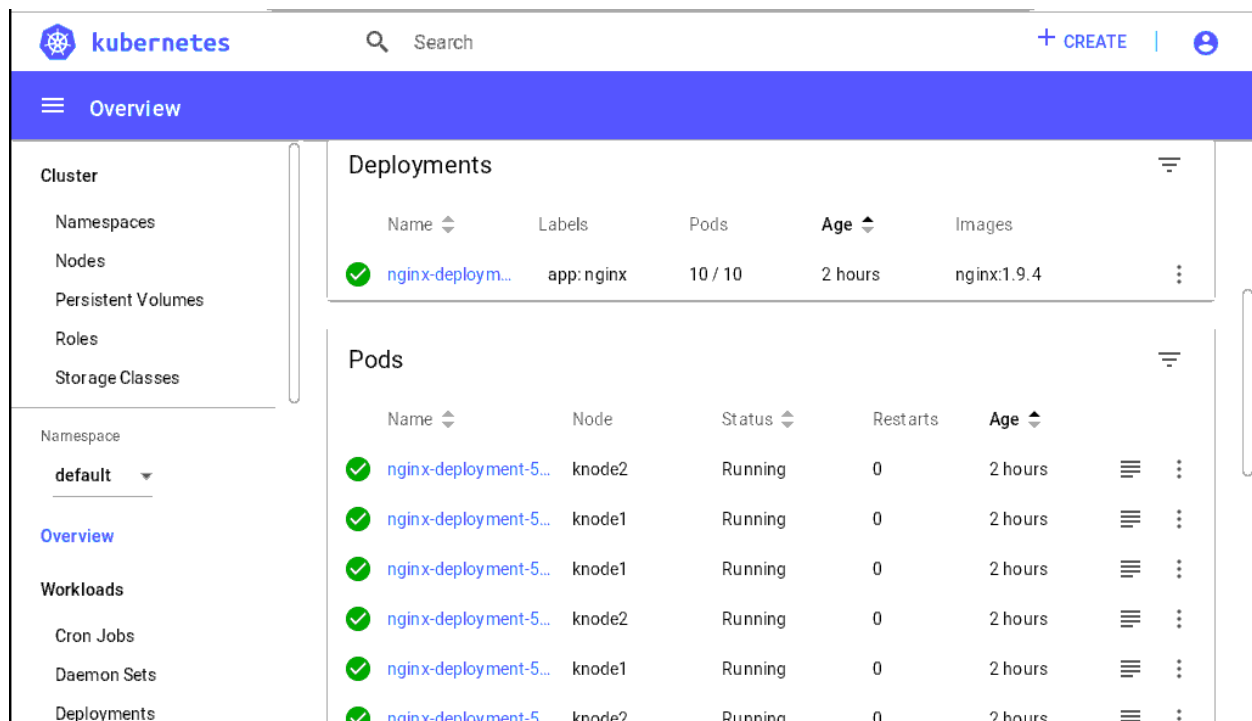
Kubernetes Dashboard overview

After authentication, dashboard will look like this

Criterion Networks



Scrolling down will show you all the deployments, which have been done till now.



Further scrolling will show you the replica sets created in previous exercises

The screenshot shows the Kubernetes dashboard interface. At the top, there's a header with the Kubernetes logo, a search bar, and a '+ CREATE' button. The sidebar on the left contains a menu with 'Overview' selected, and other options like 'Cluster', 'Namespaces', 'Nodes', 'Persistent Volumes', 'Roles', 'Storage Classes', 'Workloads', 'Cron Jobs', 'Daemon Sets', and 'Deployments'. The main area displays a table of cluster nodes and a section for Replica Sets.

Name	Labels	Pods	Age	Images
nginx-deployment-5...	app:nginx	10 / 10	2 hours	nginx:1.9.4
nginx-deployment-5...	app:nginx	0 / 0	2 hours	nginx:1.9.1
nginx-deployment-5...	app:nginx	0 / 0	2 hours	nginx:1.7.9

You can further explore the dashboard by verifying the **cluster nodes, namespaces and other side menu options** with the output of the CLI

Docker Part-III(Optional)

Exercise-31

Bring up Docker Swarm Cluster

- Install docker-ce using Exercise-11 on both dnode2 and dswarm

On dswarm1 node:

- Get the IP address of node

```
ifconfig ens3
```

```
root@dswarm1:~# ifconfig ens3
ens3      Link encap:Ethernet  HWaddr 52:54:00:9f:6a:79
          inet addr:192.168.122.133  Bcast:192.168.122.255  Mask:255.255.255.0
          inet6 addr: fe80::5054:ff:fe9f:6a79/64  Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:81842 errors:0 dropped:12 overruns:0 frame:0
          TX packets:31232 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:16530433 (16.5 MB)  TX bytes:5473350 (5.4 MB)

root@dswarm1:~#
```

- Initialize the docker swarm to advertize this IP address

```
docker swarm init --advertise-addr <docker-swarm-IP>
```

```
root@dswarm1:~# docker swarm init --advertise-addr 192.168.122.133
Swarm initialized: current node (in4u5ahm6js863qvip44b84gw) is now a manager.

To add a worker to this swarm, run the following command:

    docker swarm join --token SWMTKN-1-5oo5z5yerqdtcx9o3i8ao839yvi4ghfsrsvsuovz9ku5qd69ipq-1zh0nirbu6jzyld31p6pohljz 192.168.122.133:2377

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.

root@dswarm1:~#
```

- Observe the output displayed by above command. We will run this commands on dnode1 and dnode2 to make them part of swarm cluster

On dnode1 node

- Join the dswarm cluster

```
docker swarm join --token <<TOKEN-ID-FROM-ABOVE-OUTPUT>> <<IP-ADDRESS-OF-DSWA
RM1>:2377
```

Criterion Networks

```
root@dnode1:~# docker swarm join --token SWMTKN-1-5oo5z5yerqdtcx9o3i8ao839yvi4ghfsvsuovz9ku5qd69ipq-lzh0nirbu6jzyld31p6pohljz 192.168.122.133:2377
This node joined a swarm as a worker.
root@dnode1:~#
```

On dnode2 node

- Join the dswarm cluster

```
docker swarm join --token <<TOKEN-ID-FROM-ABOVE-OUTPUT>> <<IP-ADDRESS-OF-DSWA
RM1>:2377
```

```
root@dnode2:~# docker swarm join --token SWMTKN-1-5oo5z5yerqdtcx9o3i8ao839yvi4ghfsvsuovz9ku5qd69ipq-lzh0nirbu6jzyld31p6pohljz 192.168.122.133:2377
This node joined a swarm as a worker.
root@dnode2:~#
```

- Verify the nodes in the cluster on dswarm1 node

```
docker node ls
```

```
root@dswarm1:~# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
s9rukzpz7gxfdosh0v3llbpin	dnode1	Ready	Active		18.06.0-ce
zjazfm3uedoi66mrwzp57w6xo	dnode2	Ready	Active		18.06.0-ce
in4u5ahm6js863qvip44b84gw *	dswarm1	Ready	Active	Leader	18.06.0-ce

```
root@dswarm1:~#
```

TIP: For debugging in case of any failure, to leave the swarm cluster on a node. [DO NOT LEAVE]

docker swarm leave

Exercise-32

Deploy the App on Docker Swarm Cluster

- On "dswarm1" node, create a file named docker-compose.yml

```
vim docker-compose.yml
```

- Insert the below lines

```
version: "3"

services:

  web:

    # replace username/repo:tag with your name and image details

    image: nareshthukkani/learning:flaskapp

    deploy:

      replicas: 3

      resources:

        limits:

          cpus: "0.1"

          memory: 50M

      restart_policy:

        condition: on-failure

    ports:

      - "5000:80"

    networks:

      - webnet

networks:

  webnet:
```

```
version: "3"
services:
  web:
    # replace username/repo:tag with your name and image details
    image: nareshthukkani/learning:flaskapp
    deploy:
      replicas: 3
      resources:
        limits:
          cpus: "0.1"
          memory: 50M
      restart_policy:
        condition: on-failure
    ports:
      - "5000:80"
    networks:
      - webnet
networks:
  webnet:
```

- Deploy the stack

```
docker stack deploy -c docker-compose.yml getstartedlab
```

```
root@dswarm1:~# docker stack deploy -c docker-compose.yml getstartedlab
Creating network getstartedlab_webnet
Creating service getstartedlab_web
root@dswarm1:~#
```

- Verify multiple containers are deployed as per docker-compose.yml

```
docker stack ps getstartedlab
```

```
root@dswarm1:~# docker stack ps getstartedlab
```

ID	PORTS	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR
inp4te3gr1o0		getstartedlab_web.1	nareshthukkani/learning:flaskapp	dnode2	Running	Running 10 seconds ago	
ij8oxu8aqhk8		getstartedlab_web.2	nareshthukkani/learning:flaskapp	dswarm1	Running	Running 9 seconds ago	
9b91j8w2luqa		getstartedlab_web.3	nareshthukkani/learning:flaskapp	dnode1	Running	Running 17 seconds ago	

```
root@dswarm1:~#
```

- Verify using VNC, Web browser by pointing at following IP:port

```
http://<IP-dnode1>:5000
```

```
http://<IP-dnode2>:5000
```

```
http://<IP-dswarm>:5000
```

- For clients that are OUTSIDE the cluster and accessing the EXPOSED services, Docker swarm uses Routing Mesh to ensure services are reachable from any node. Verify IP Tables rules on all nodes

```
iptables -t nat -L
```

```
root@dswarm1:~# iptables -t nat -L
Chain PREROUTING (policy ACCEPT)
target     prot opt source                destination
DOCKER-INGRESS  all  --  anywhere                anywhere
DOCKER      all  --  anywhere                anywhere
ADDRTYPE match dst-type LOCAL
ADDRTYPE match dst-type LOCAL

Chain INPUT (policy ACCEPT)
target     prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination
DOCKER-INGRESS  all  --  anywhere                anywhere
DOCKER      all  --  anywhere                !127.0.0.0/8
ADDRTYPE match dst-type LOCAL
ADDRTYPE match dst-type LOCAL

Chain POSTROUTING (policy ACCEPT)
target     prot opt source                destination
MASQUERADE  all  --  anywhere                anywhere
MASQUERADE  all  --  172.18.0.0/16            anywhere
MASQUERADE  all  --  172.17.0.0/16            anywhere
ADDRTYPE match src-type LOCAL

Chain DOCKER (2 references)
target     prot opt source                destination
RETURN     all  --  anywhere                anywhere
RETURN     all  --  anywhere                anywhere

Chain DOCKER-INGRESS (2 references)
target     prot opt source                destination
DNAT       tcp  --  anywhere                anywhere
RETURN     all  --  anywhere                anywhere
tcp dpt:5000 to:172.18.0.2:5000

root@dswarm1:~#
```

- Docker swarm creates hidden containers that does IPVS Linux Kernel based load-balancing.
- On any nodes where service is present, check the network namespaces created.

```
cd /var/run/docker/netns
ls
```

```
root@dnode1:~# cd /var/run/docker/netns
root@dnode1:/var/run/docker/netns# ls
1-fxbsfdbj84 1-mirs077ko3 63f98531dc55 6c251bb6d3f4 88a0e4b2cb48 c4fbc400ed7d ea4a14ccd4e3 ec419c814570 ingress_sbox
root@dnode1:/var/run/docker/netns#
```

(You will see a network namespace with name ingress_sbox)

- Enter the superuser mode and then execute the command

```
sudo su
nsenter --net=ingress_sbox /bin/bash
iptables -t mangle -L
```

```

root@dnode1:/var/run/docker/netns# sudo su
root@dnode1:/run/docker/netns# nsenter --net=ingress_sbox /bin/bash
root@dnode1:/run/docker/netns# iptables -t mangle -L
Chain PREROUTING (policy ACCEPT)
target     prot opt source                destination              tcp dpt:5000 MARK set 0x103

Chain INPUT (policy ACCEPT)
target     prot opt source                destination              MARK set 0x103

Chain FORWARD (policy ACCEPT)
target     prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination

Chain POSTROUTING (policy ACCEPT)
target     prot opt source                destination
root@dnode1:/run/docker/netns#

```

(Verify the MARK for published service port)

- Verify the IPVS Load Balancing rules in Round Robin Fashion

```

apt-get install ipvsadm

ipvsadm -L

```

```

root@dnode1:/run/docker/netns# ipvsadm -L
IP Virtual Server version 1.2.1 (size=4096)
Prot LocalAddress:Port Scheduler Flags
  -> RemoteAddress:Port          Forward Weight ActiveConn InActConn
FWM  259 rr
  -> 10.255.0.6:0                Masq    1      0      0
  -> 10.255.0.7:0                Masq    1      0      0
  -> 10.255.0.8:0                Masq    1      0      0
root@dnode1:/run/docker/netns#

```

- For clients that are INSIDE the cluster and accessing other INSIDE services, docker swarm provisions "IPTables" rules and IPVS Load-balancing rules on every container Namespace.
- In order to know the container's network namespace, grep the "SandboxID" from the docker inspect command. Observe the first 12 characters of the SandboxID. This will be the Network namespace ID.

```

docker ps

docker inspect CONTAINERID | grep -i SandboxID

```

- Enter the superuser mode and then execute the command

```

sudo su

nsenter --net=CONTAINER-NETNS-ID /bin/bash

```

```
iptables -t mangle -L
```

(Verify the MARK for published service port)

- Verify the IPVS Load Balancing rules in Round Robin Fashion

```
ipvsadm -L
```

CONFIDENTIAL

Exercise-33

Docker Stacks with Visualizer

On dswarm1 node

- Create new file that includes a visualizer along with the contents in docker-compose.yml

```
vim docker-compose-1.yml
```

```
root@dswarm1:~# vim docker-compose-1.yml
root@dswarm1:~#
```

- Insert the below lines

```
version: "3"

services:

  web:

    # replace username/repo:tag with your name and image details

    image: nareshthukkani/learning:flaskapp

    deploy:

      replicas: 6

      restart_policy:

        condition: on-failure

      resources:

        limits:

          cpus: "0.1"

          memory: 50M

    ports:

      - "80:80"

    networks:
```

```
- webnet

visualizer:

  image: dockersamples/visualizer:stable

  ports:

    - "8080:8080"

  volumes:

    - "/var/run/docker.sock:/var/run/docker.sock"

  deploy:

    placement:

      constraints: [node.role == manager]

    networks:

      - webnet

networks:

  webnet:
```

```

version: "3"
services:
  web:
    # replace username/repo:tag with your name and image details
    image: nareshthukkani/learning:flaskapp
    deploy:
      replicas: 6
      restart_policy:
        condition: on-failure
      resources:
        limits:
          cpus: "0.1"
          memory: 50M
    ports:
      - "80:80"
    networks:
      - webnet
  visualizer:
    image: dockersamples/visualizer:stable
    ports:
      - "8080:8080"
    volumes:
      - "/var/run/docker.sock:/var/run/docker.sock"
    deploy:
      placement:
        constraints: [node.role == manager]
    networks:
      - webnet
networks:
  webnet:
~
~

```

- Deploy the stack

```
docker stack deploy -c docker-compose-1.yml getstartedlab
```

```

root@dswarm1:~# docker stack deploy -c docker-compose-1.yml getstartedlab
Updating service getstartedlab_web (id: 538hvown2eplyfbfxwvijhn8a)
Creating service getstartedlab_visualizer
root@dswarm1:~#

```

- Verify visualizer using VNC, Web browser at port [IP-dswarm:5000]
- Visualize the docker containers across multiple nodes

Exercise-34

Docker stacks with Persistent Data

On dswarm1 node

- Create a new file that includes Redis database service along with contents of docker-compose-1.yml

```
vim docker-compose-2.yml
```

```
root@dswarm1:~# vim docker-compose-2.yml
```

- Insert the below lines

```
version: "3"

services:

  web:

    # replace username/repo:tag with your name and image details
    image: nareshthukkani/learning:flaskapp

    deploy:

      replicas: 6

      restart_policy:

        condition: on-failure

    resources:

      limits:

        cpus: "0.1"

        memory: 50M

    ports:

      - "80:80"

    networks:
```

```
- webnet

visualizer:

  image: dockersamples/visualizer:stable

  ports:

    - "8080:8080"

  volumes:

    - "/var/run/docker.sock:/var/run/docker.sock"

  deploy:

    placement:

      constraints: [node.role == manager]

    networks:

      - webnet

redis:

  image: redis

  ports:

    - "6379:6379"

  volumes:

    - "/root/learning/data:/data"

  deploy:

    placement:

      constraints: [node.role == manager]

  command: redis-server --appendonly yes

  networks:

    - webnet
```

```
networks:
```

```
  webnet:
```

```
version: "3"
services:
  web:
    # replace username/repo:tag with your name and image details
    image: nareshthukkani/learning:flaskapp
    deploy:
      replicas: 6
      restart_policy:
        condition: on-failure
      resources:
        limits:
          cpus: "0.1"
          memory: 50M
    ports:
      - "80:80"
    networks:
      - webnet
  visualizer:
    image: dockersamples/visualizer:stable
    ports:
      - "8080:8080"
    volumes:
      - "/var/run/docker.sock:/var/run/docker.sock"
    deploy:
      placement:
        constraints: [node.role == manager]
    networks:
      - webnet
  redis:
    image: redis
    ports:
      - "6379:6379"
    volumes:
      - "/root/learning/data:/data"
    deploy:
```

- Create a directory for persistent data under /root/learning path

```
mkdir /root/learning
cd /root/learning
mkdir data
```

- Deploy the stack

```
docker stack deploy -c /home/ubuntu/docker-compose-2.yml getstartedlab
```

```
root@dswarm1:/root/learning# docker stack deploy -c /home/ubuntu/docker-compose-2.yml getstartedlab
Updating service getstartedlab_web (id: 538hvown2eplyfbfxwvijn8a)
Updating service getstartedlab_visualizer (id: ofyymfqyyuk8r30h1miq397mx)
Creating service getstartedlab_redis
root@dswarm1:/root/learning#
```

- Verify visualizer using VNC, Web browser at port [IP-dswarm:8080]
- Visualize the docker containers across multiple nodes
- Visualize Redis service is present on dswarm Manager only
- Even though the Redis container terminates, data is persistent in /root/learning/data folder

Miscellaneous

Track Labels

Track labels

Guidance: You will need to create files based on inputs provided here

- The primary, stable release would have a track label with value as stable:

```
name: frontend

  replicas: 3

  ...

  labels:

    app: guestbook

    tier: frontend

    track: stable

  ...

  image: gb-frontend:v3
```

and then you can create a new release of the guestbook frontend that carries the track label with different value (i.e. canary), so that two sets of pods would not overlap:

```
name: frontend-canary

replicas: 1

...

labels:

  app: guestbook

  tier: frontend

  track: canary

...
```



```
image: gb-frontend:v4
```

The frontend service would span both sets of replicas by selecting the common subset of their labels (i.e. omitting the track label), so that the traffic will be redirected to both applications:

```
selector:

  app: guestbook

  tier: frontend
```

You can tweak the number of replicas of the stable and canary releases to determine the ratio of each release that will receive live production traffic (in this case, 3:1). Once you're confident, you can update the stable track to the new application release and remove the canary one.

Updating the Labels and annotations

Guidance: Explore below commands and play with deployments

Sometimes existing pods and other resources need to be relabeled before creating new resources. This can be done with `kubectl label`. For example, if you want to label all your `nginx` pods as frontend tier, simply run:

```
kubectl label pods -l app=nginx tier=fe
```

Output:

```
pod "my-nginx-2035384211-j5fhi" labeled
pod "my-nginx-2035384211-u2c7e" labeled
pod "my-nginx-2035384211-u3t6x" labeled
```

This first filters all pods with the label `app=nginx`, and then labels them with the `tier=fe`. To see the pods you just labeled, run:

```
kubectl get pods -l app=nginx -L tier
```

Output:

NAME	READY	STATUS	RESTARTS	AGE	TIER
my-nginx-2035384211-j5fhi	1/1	Running	0	23m	fe
my-nginx-2035384211-u2c7e	1/1	Running	0	23m	fe

my-nginx-2035384211-u3t6x	1/1	Running	0	23m	fe
---------------------------	-----	---------	---	-----	----

This outputs all app=nginx pods, with an additional label column of pods tier (specified with -L or --label-columns).

Updating annotations

Sometimes you would want to attach annotations to resources. Annotations are arbitrary non-identifying metadata for retrieval by API clients such as tools, libraries, etc. This can be done with `kubectl annotate`. For example:

```
kubectl annotate pods my-nginx-v4-9gw19 description='my frontend running nginx'

kubectl get pods my-nginx-v4-9gw19 -o yaml
```

Output:

```
apiversion: v1

kind: pod

metadata:

  annotations:

    description: my frontend running nginx

...
```